

# Lecture 2: Uninformed Search

Shuai Li

John Hopcroft Center, Shanghai Jiao Tong University


<https://shuaili8.github.io>

<https://shuaili8.github.io/Teaching/CS410/index.html>



Part of slide credits: CMU AI & <http://ai.berkeley.edu>

# Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search




# Agents that Plan



# Rationality


- What is rational depends on:
  - Performance measure
  - Agent's prior knowledge of environment
  - Actions available to agent
  - Percept/sensor sequence to date
- Being rational means **maximizing your expected utility**




# Rational Agents

- Are rational agents *omniscient*? 无所不知的
  - No – they are limited by the available percepts
- Are rational agents *clairvoyant*? 透视的
  - No – they may lack knowledge of the environment dynamics
- Do rational agents *explore* and *learn*?
  - Yes – in unknown environments these are essential
- So rational agents are not necessarily successful, but they are *autonomous* (i.e., control their own behavior)

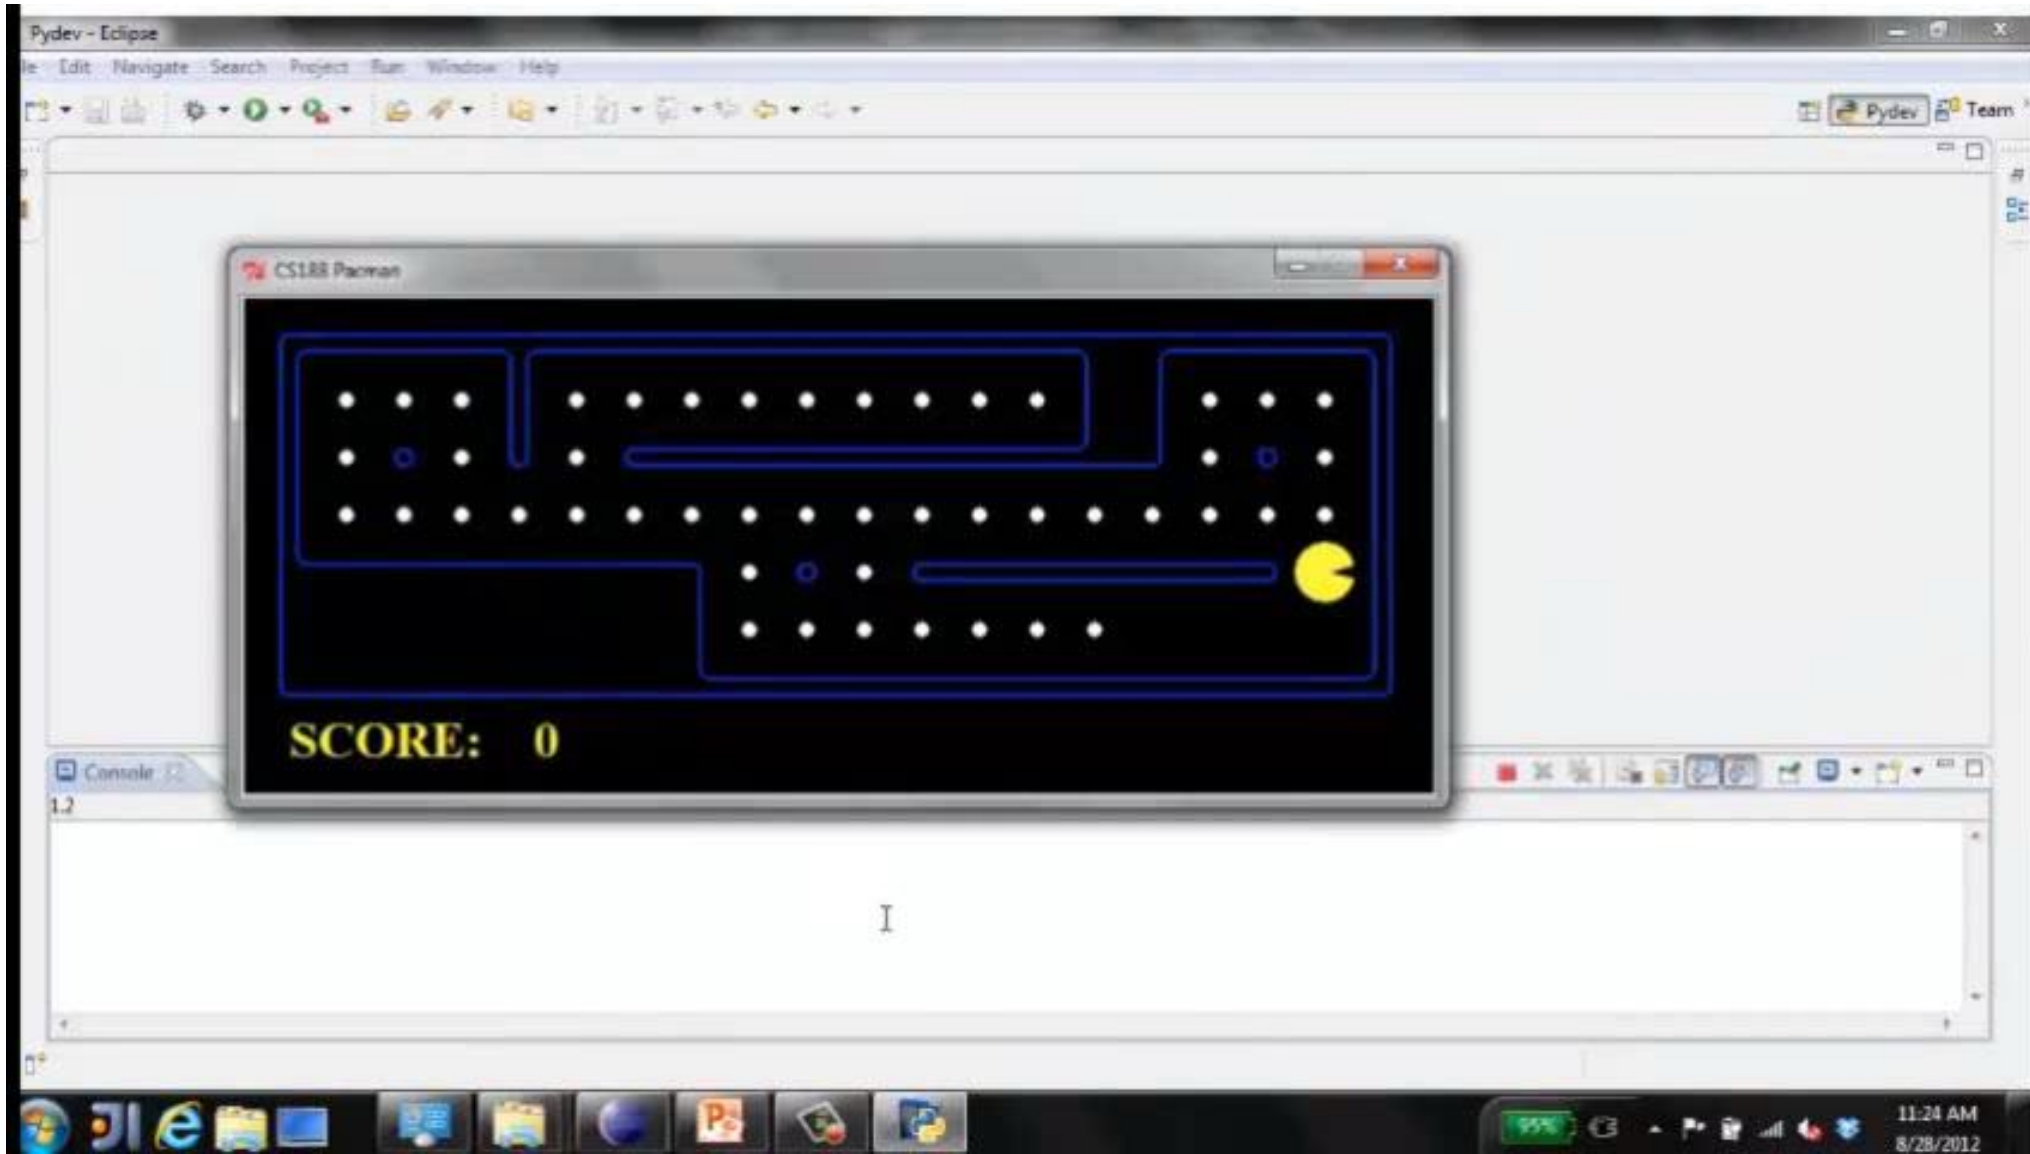
# Planning Agents



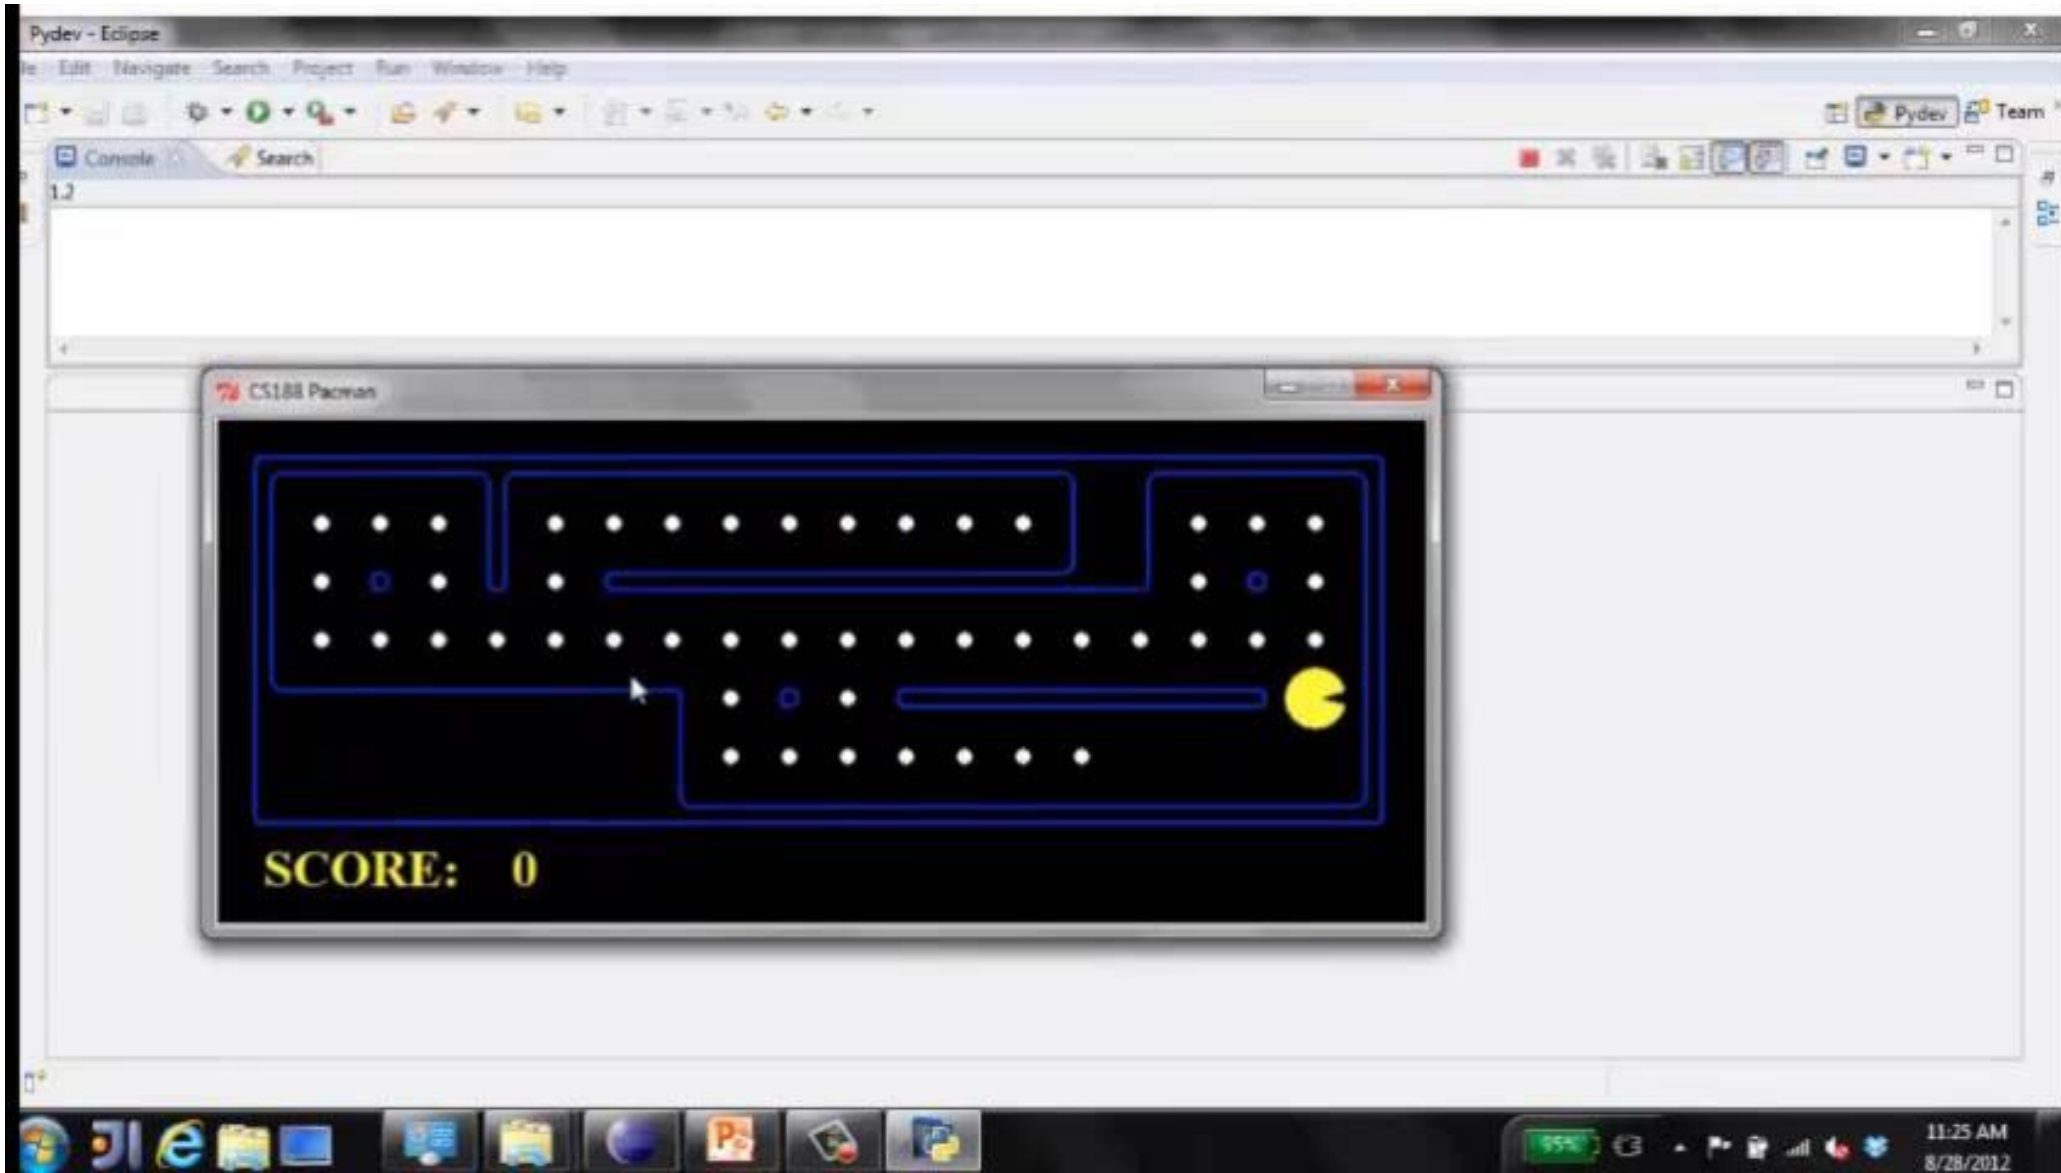
- Planning agents:
  - Ask “what if”
  - Decisions based on (hypothesized or **predicted**) consequences of actions
  - Must have a **transition** model of how the world evolves in response to actions
  - Must formulate a goal (test)
  - **Consider how the world WOULD BE**
- Spectrum of deliberativeness:
  - Generate complete, optimal plan offline, then execute
  - Generate a simple, greedy plan, start executing, replan when something goes wrong
- Optimal vs. complete planning
- Planning vs. replanning
  - [Demo: re-planning (L2D3)]
  - [Demo: mastermind (L2D4)]



# Video of Demo Replanning



# Video of Demo Mastermind






# Search Problems




# Search Problems

- A **search problem** consists of:

- A state space




- For each state, a set **Actions(s)** of successors/actions



{N, E}

- A transition model  $T(s,a)$

- A step cost(reward) function  $c(s,a,s')$



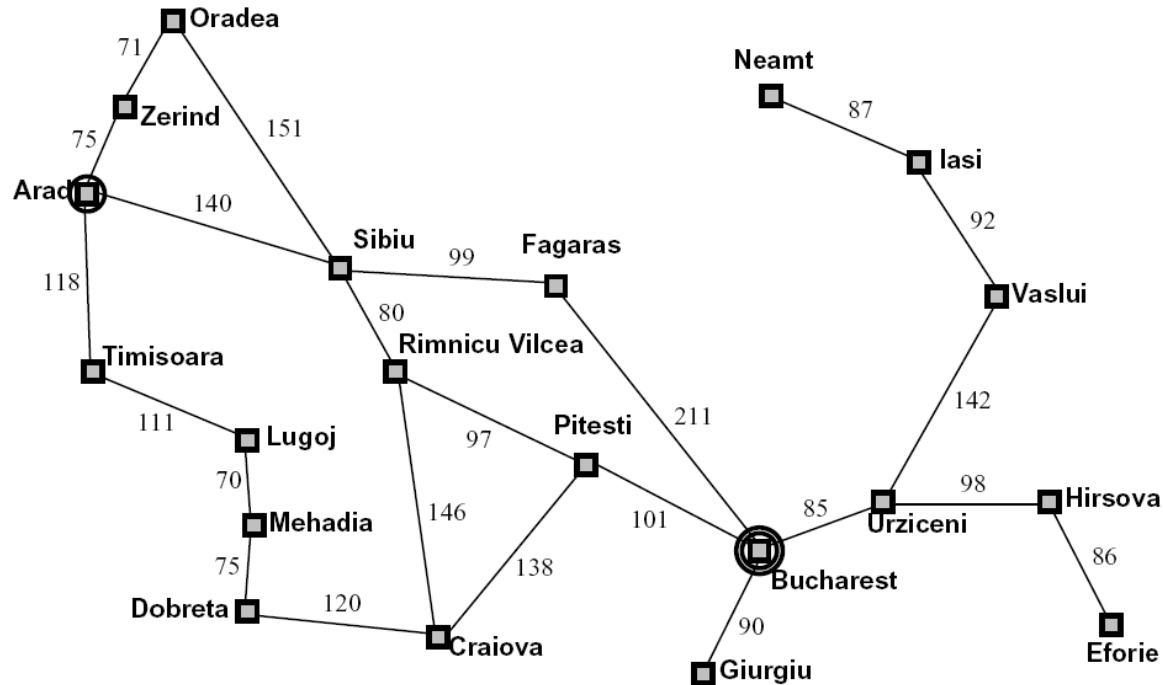
- A start state and a goal test

- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

# Search Problems Are Models




# Example: Traveling in Romania



- State space:
  - Cities
- Successor function:
  - Roads: Go to adjacent city with cost = distance
- Start state:
  - Arad
- Goal test:
  - Is state == Bucharest?
- Solution?

# What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

- **Problem: Pathing**


- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is (x,y)=END

- **Problem: Eat-All-Dots**


- States: {(x,y), dot booleans}
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

# State Space Sizes?

- World state:
  - Agent positions: 120
  - Food count: 30
  - Ghost positions: 12
  - Agent facing: NSEW
- How many
  - World states?  
 $120 \times (2^{30}) \times (12^2) \times 4$
  - States for pathing?  
120
  - States for eat-all-dots?  
 $120 \times (2^{30})$




# Safe Passage



- Problem: eat all dots while keeping the ghosts perma-scared
- What does the state space have to specify?
  - (agent position, dot booleans, power pellet booleans, remaining scared time)


# State Space Graphs and Search Trees






# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



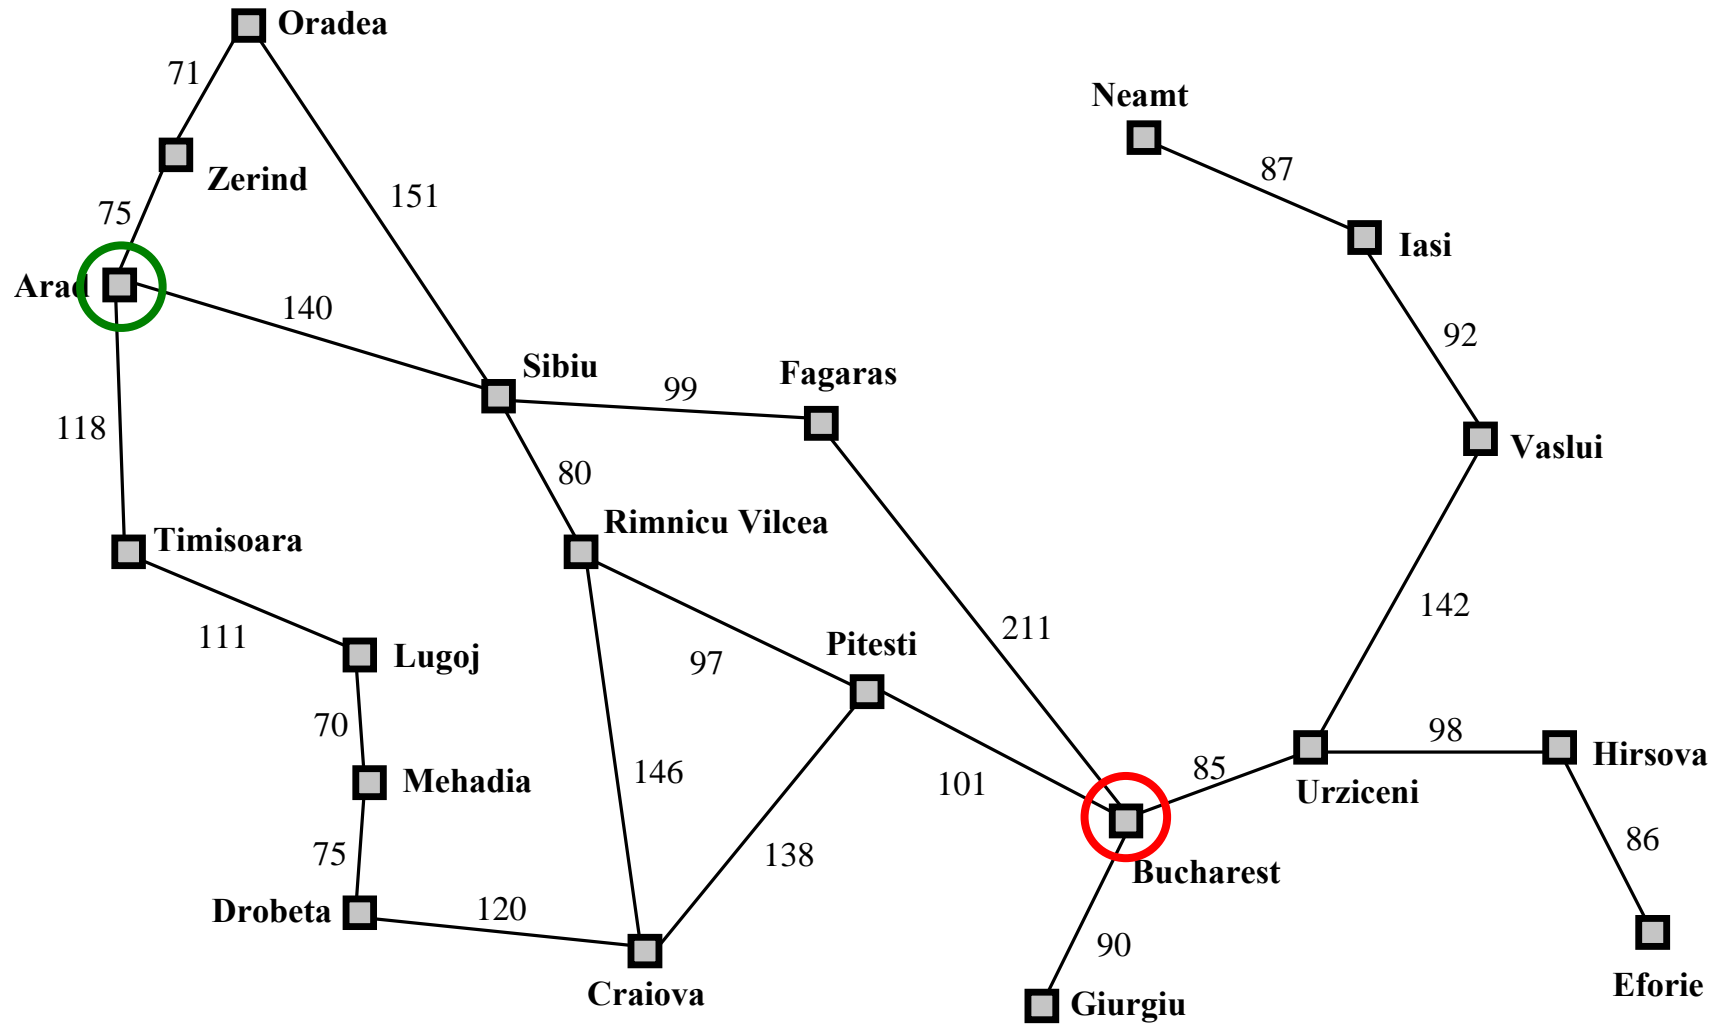
# State Space Graphs

- State space graph: A mathematical representation of a search problem
  - Nodes are (abstracted) world configurations
  - Arcs represent successors (action results)
  - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




*Tiny search graph for a tiny search problem*


# More Examples



# More Examples




# Search Trees



- A search tree:
  - A “what if” tree of plans and their outcomes
  - The start state is the root node
  - Children correspond to successors
  - Nodes show states, but correspond to **PLANS** that achieve those states
  - For most problems, we can never actually build the whole tree

# State Space Graphs vs. Search Trees


## State Space Graph



*Each NODE in in the search tree is an entire PATH in the state space graph.*


*We construct both on demand – and we construct as little as possible.*

## Search Tree




# State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?



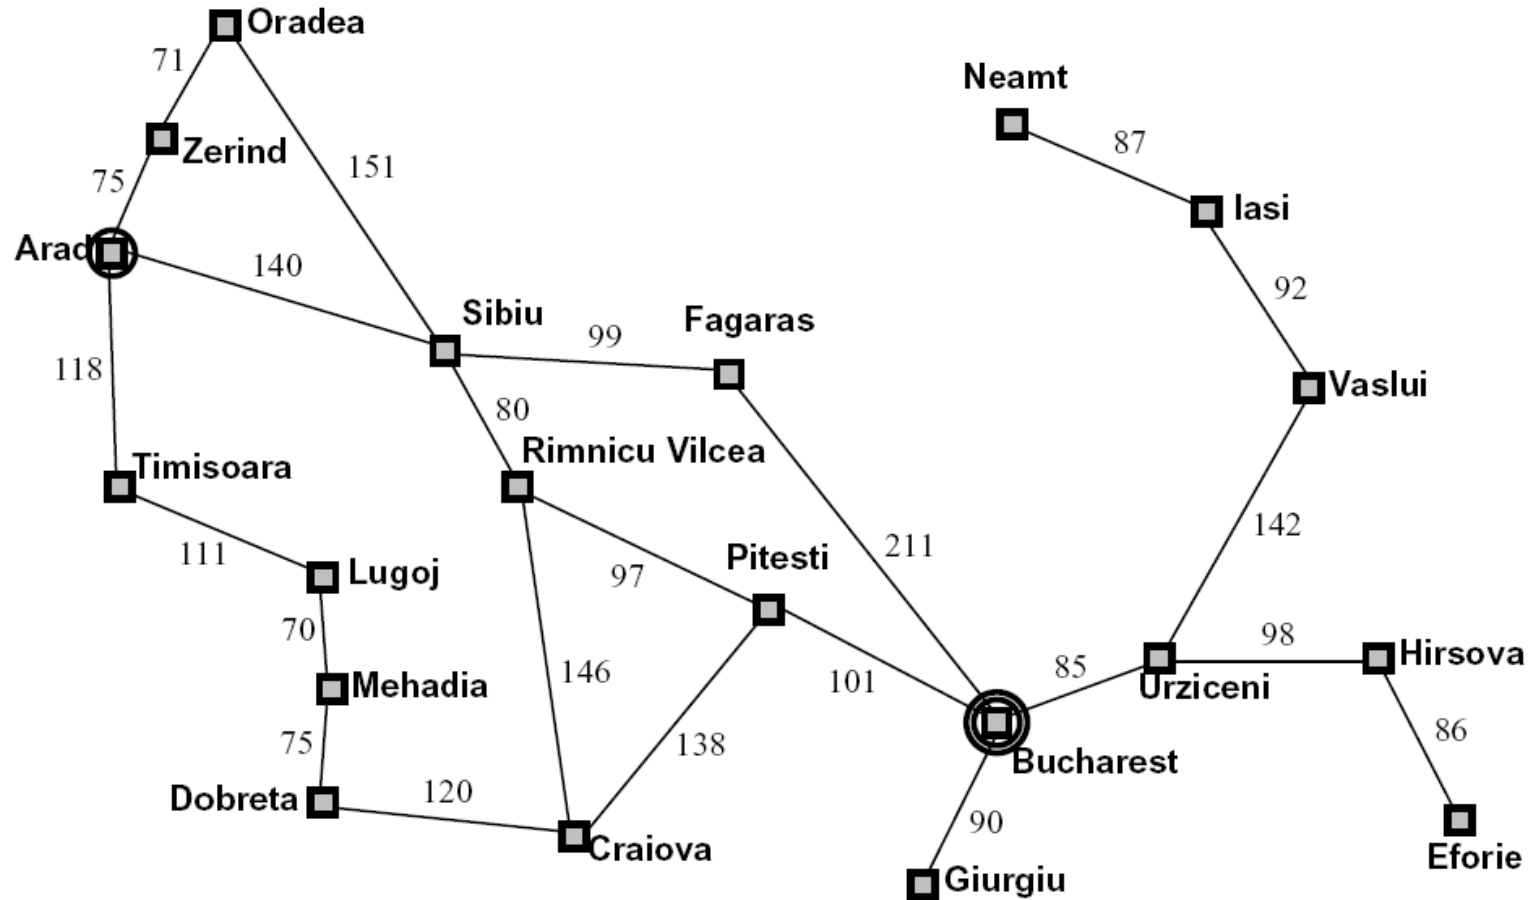
Important: Lots of repeated structure in the search tree!

# Tree Search






# Search Example: Romania



# Searching with a Search Tree



- Search:
  - Expand out potential plans (tree nodes)
  - Maintain a **fringe** of partial plans under consideration
  - Try to expand as few tree nodes as possible

# General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```

- Important ideas:
  - Fringe
  - Expansion
  - Exploration strategy
- Main question: which fringe nodes to explore?

# General Tree Search 2

function TREE\_SEARCH(problem) returns a solution, or failure

initialize the frontier as a specific work list (stack, queue, priority queue)

add initial state of problem to frontier

loop do

if the frontier is empty then

return failure

**choose** a node and remove it from the frontier



if the node contains a goal state then

return the corresponding solution

for each resulting child from node

add child to the frontier

# Example: Tree Search



~~s~~

~~s → d~~

s → e

s → p

s → d → b

s → d → c

~~s → d → e~~

s → d → e → h

~~s → d → e → r~~

~~s → d → e → r → f~~

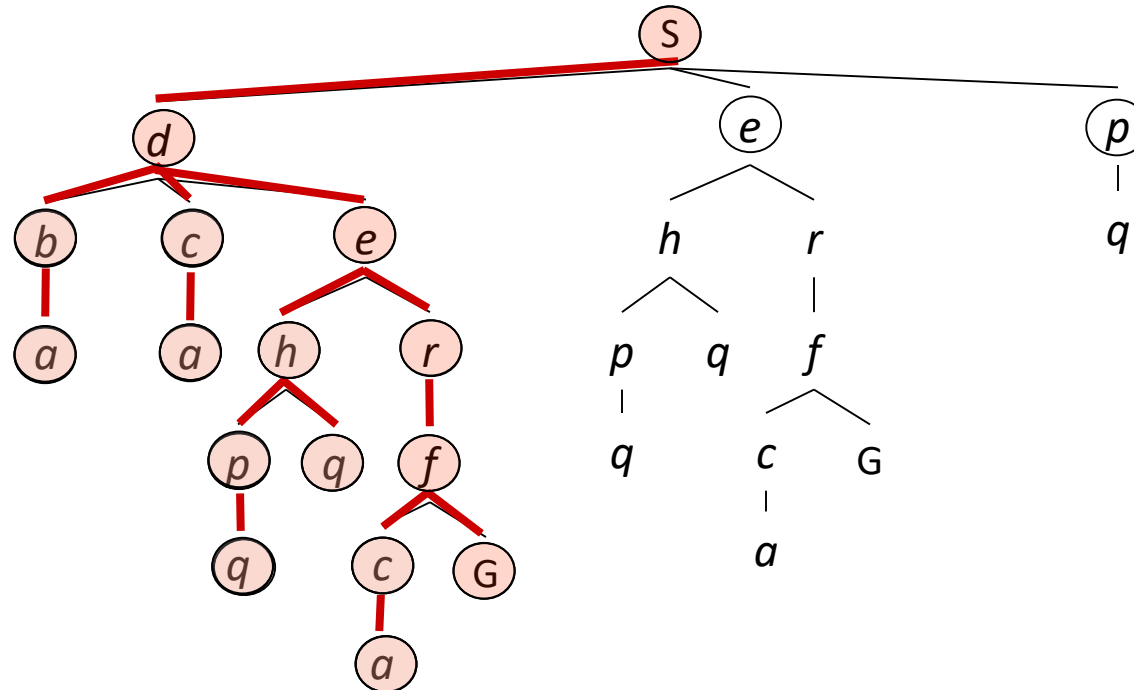

s → d → e → r → f → c

~~s → d → e → r → f → G~~

# Depth-First (Tree) Search

Strategy: expand a  
deepest node first



Implementation:  
Fringe is a *LIFO* stack



# Breadth-First (Tree) Search



Strategy: expand a  
shallowest node first

Implementation: Fringe  
is a *FIFO queue*



# Search Algorithm Properties


- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?
- Cartoon of search tree:
  - $b$  is the branching factor
  - $m$  is the maximum depth
  - solutions at various depths
- Number of nodes in entire tree?
  - $1 + b + b^2 + \dots + b^m = O(b^m)$





# Depth-First Search (DFS) Properties

- What nodes DFS expand?
  - Some left prefix of the tree.
  - Could process the whole tree!
  - If  $m$  is finite, takes time  $O(b^m)$
- How much space does the fringe take?
  - Only has siblings on path to root, so  $O(bm)$
- Is it complete?
  - $m$  could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
  - No, it finds the “leftmost” solution, regardless of depth or cost



# Breadth-First Search (BFS) Properties

- What nodes does BFS expand?

- Processes all nodes above shallowest solution
- Let depth of shallowest solution be  $s$
- Search takes time  $O(b^s)$

$s$  tiers

- How much space does the fringe take?


- Has roughly the last tier, so  $O(b^s)$

- Is it complete?


- $s$  must be finite if a solution exists

- Is it optimal?

- Only if costs are all 1 (more on costs later)




# DFS vs BFS




# DFS vs BFS

- When will BFS outperform DFS?
- When will DFS outperform BFS?

# Video of Demo Maze Water DFS/BFS (part 1)




# Video of Demo Maze Water DFS/BFS (part 2)




# Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
  - Run a DFS with depth limit 1. If no solution...
  - Run a DFS with depth limit 2. If no solution...
  - Run a DFS with depth limit 3. ....
- Isn't that wastefully redundant?
  - Generally most work happens in the lowest level searched, so not so bad!




# A Note on Implementation

- Nodes have
  - `state`, `parent`, `action`, `path-cost`
- A child of `node` by action `a` has
  - `state` = `Transition(node.state, a)`
  - `parent` = `node`
  - `action` = `a`
  - `path-cost` = `node.path_cost` + `step_cost(node.state, a, self.state)`
- Extract solution by tracing back parent pointers, collecting actions



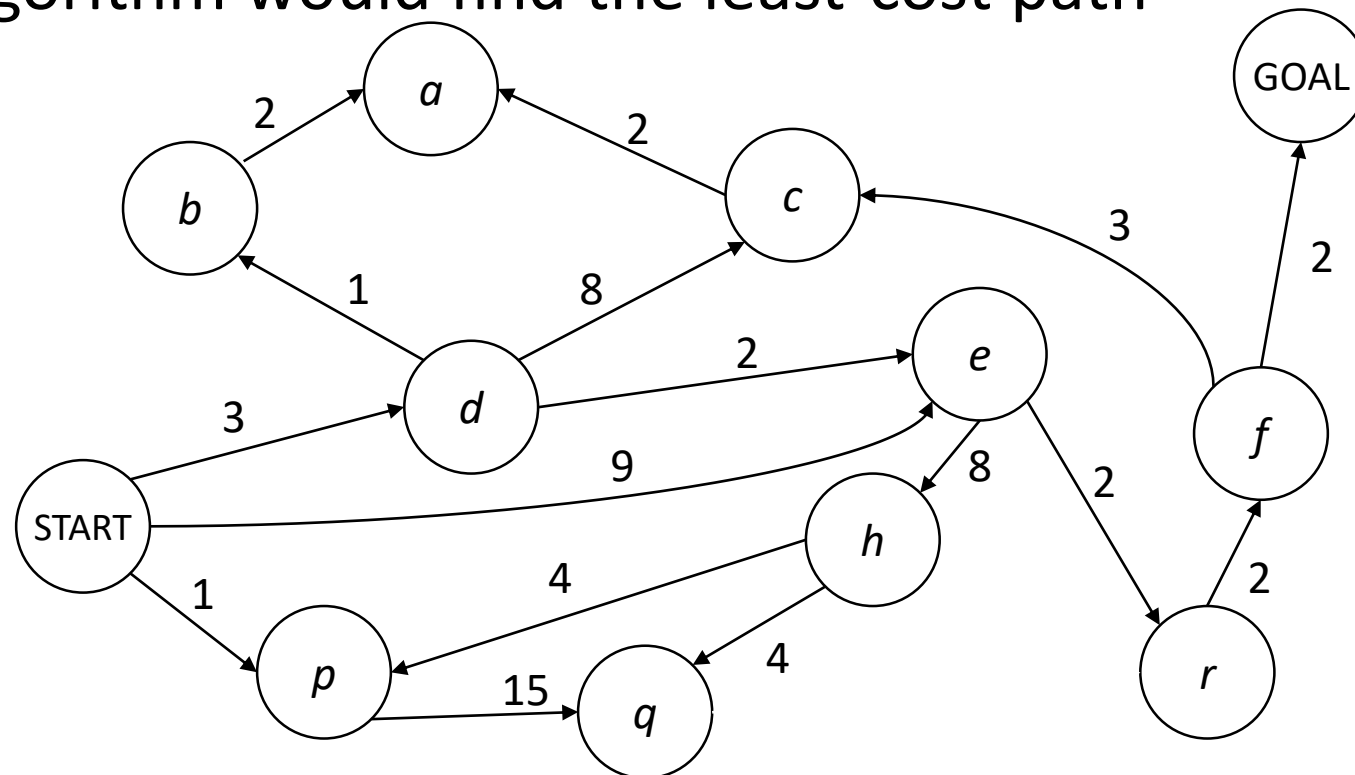


# Uniform Cost Search



# Finding a Least-Cost Path

- BFS finds the shortest path in terms of number of actions, but not the least-cost path
- A similar algorithm would find the least-cost path





How?

# Uniform Cost Search

Strategy: expand a cheapest node first:

Fringe is a priority queue  
(priority: *cumulative cost*)



# Uniform Cost Search 2

**function** UNIFORM-COST-SEARCH(**problem**) **returns** a solution, or failure

initialize the **frontier** as a **priority queue** using node's **path\_cost** as the **priority**

add initial state of **problem** to **frontier** with **path\_cost** = 0

**loop do**

**if** the **frontier** is empty **then**

**return** failure


    choose a **node** (with minimal **path\_cost**) and remove it from the **frontier**

**if** the **node** contains a goal state **then**


**return** the corresponding solution

    for each resulting **child** from node


        add **child** to the **frontier** with **path\_cost** = **path\_cost**(node) + **cost**(node, **child**)



# Walk-through UCS




# Walk-through UCS





# Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
  - Processes all nodes with cost less than cheapest solution!
  - If that solution costs  $C^*$  and arcs cost at least  $\varepsilon$ , then the “effective depth” is roughly  $C^*/\varepsilon$
  - Takes time  $O(b^{C^*/\varepsilon})$  (exponential in effective depth)
- How much space does the fringe take?
  - Has roughly the last tier, so  $O(b^{C^*/\varepsilon})$
- Is it complete?
  - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
  - Yes! (Proof next via  $A^*$ )



# Uniform Cost Issues


- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
  - Explores options in every “direction”
  - No information about goal location
- We’ll fix that soon!




[Demo: empty grid UCS (L2D5)]  
[Demo: maze with deep/shallow  
water DFS/BFS/UCS (L2D7)]<sup>48</sup>




# Video of Demo Empty UCS (same cost)




# DFS, BFS, or UCS?




# Video of Demo Maze with Deep/Shallow Water (part 1)



# Video of Demo Maze with Deep/Shallow Water (part 2)




# Video of Demo Maze with Deep/Shallow Water (part 3)




# The One Queue

- All these search algorithms are the same except for fringe strategies
  - Conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities)
  - Practically, for DFS and BFS, you can avoid the  $\log(n)$  overhead from an actual priority queue, by using stacks and queues
  - Can even code one implementation that takes a variable queuing object






# Search and Models

- Search operates over models of the world
  - The agent doesn't actually try all the plans out in the real world!
  - Planning is all “in simulation”
  - Your search is only as good as your models...



# Search Gone Wrong?





# Summary

- Rational agents
- Search problems
- Uninformed Search Methods
  - Depth-First Search
  - Breadth-First Search
  - Uniform-Cost Search

**Shuai Li**

<https://shuaili8.github.io>

## Questions?