

# Lecture 8: Recurrent Neural Networks

Shuai Li

John Hopcroft Center, Shanghai Jiao Tong University

<https://shuaili8.github.io>

<https://shuaili8.github.io/Teaching/VE445/index.html>




# Motivation

- Not all problems can be converted into one with fixed length inputs and outputs
- Problems such as *Speech Recognition* or *Time-series Prediction* require a system to store and use context information
  - Simple case:
    - Output YES if the number of 1s is even, else NO!  
1000010101 – YES, 100011 – NO, ...
- Hard/Impossible to choose a fixed context window
  - There can always be a new sample longer than anything seen

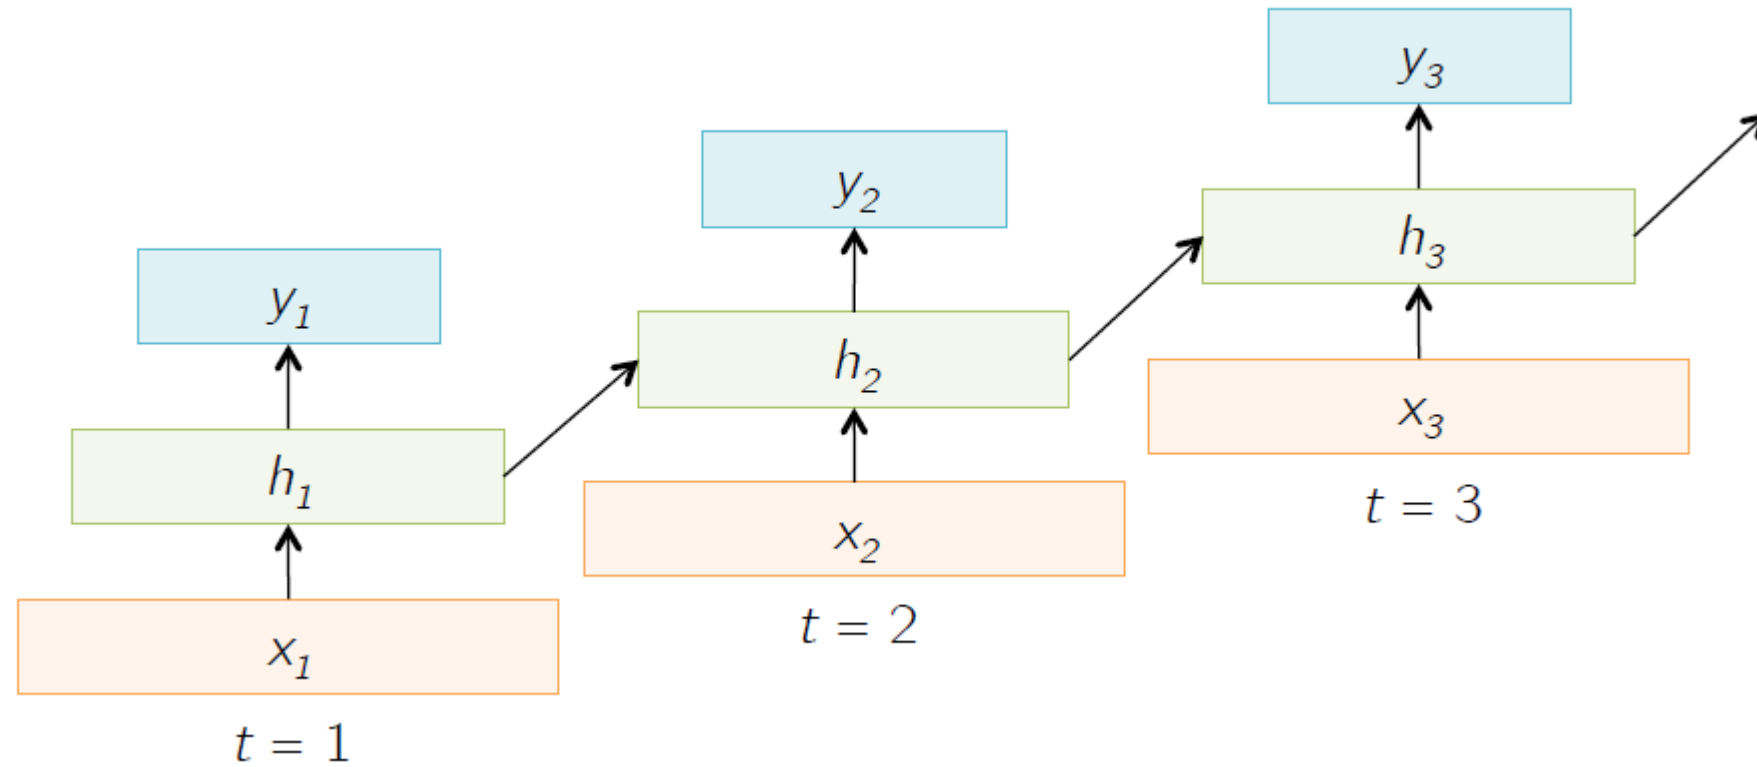
# Recurrent neural networks (RNNs)

- Recurrent neural networks take the previous output or hidden states as inputs. The composite input at time  $t$  has some historical information about the happenings at time  $T < t$
- RNNs are useful as their intermediate values (state) can store information about past inputs for a time that is not fixed a priori

# Feed-forward Network




# RNN example



# The vanilla RNN cell

- The state consists of a single “*hidden*” vector  $\mathbf{h}$ :




$$h_t = f_W(h_{t-1}, x_t)$$




$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$

# Computational graph of RNN

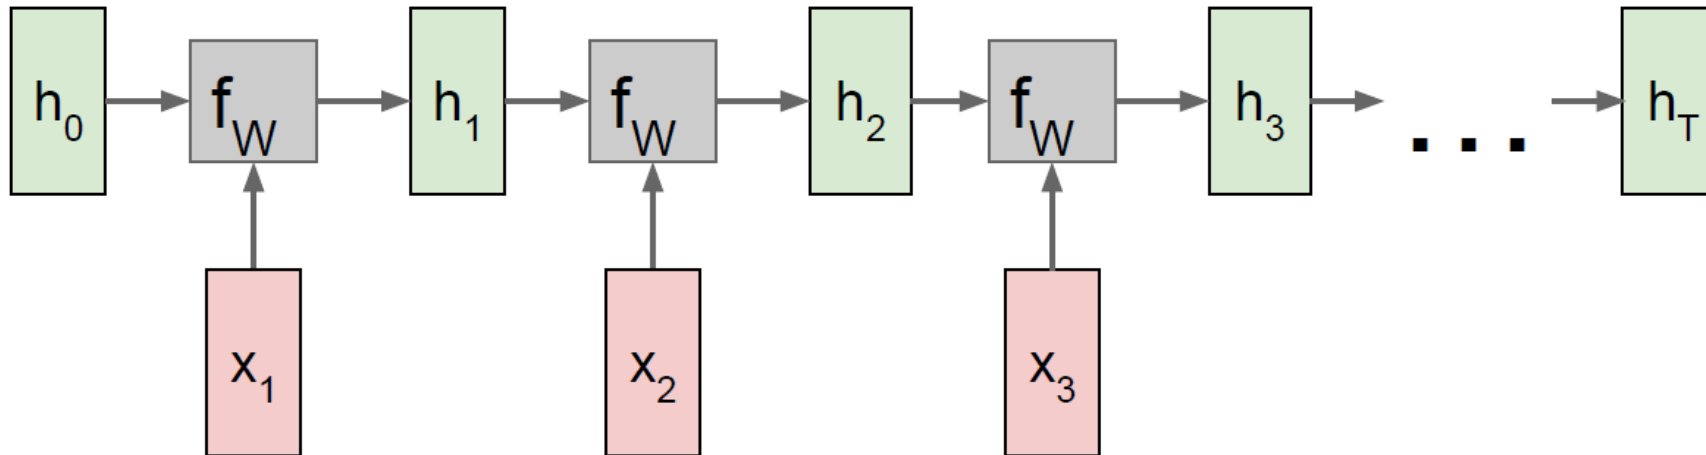


# Computational graph of RNN






# Computational graph of RNN




# Computational graph of RNN

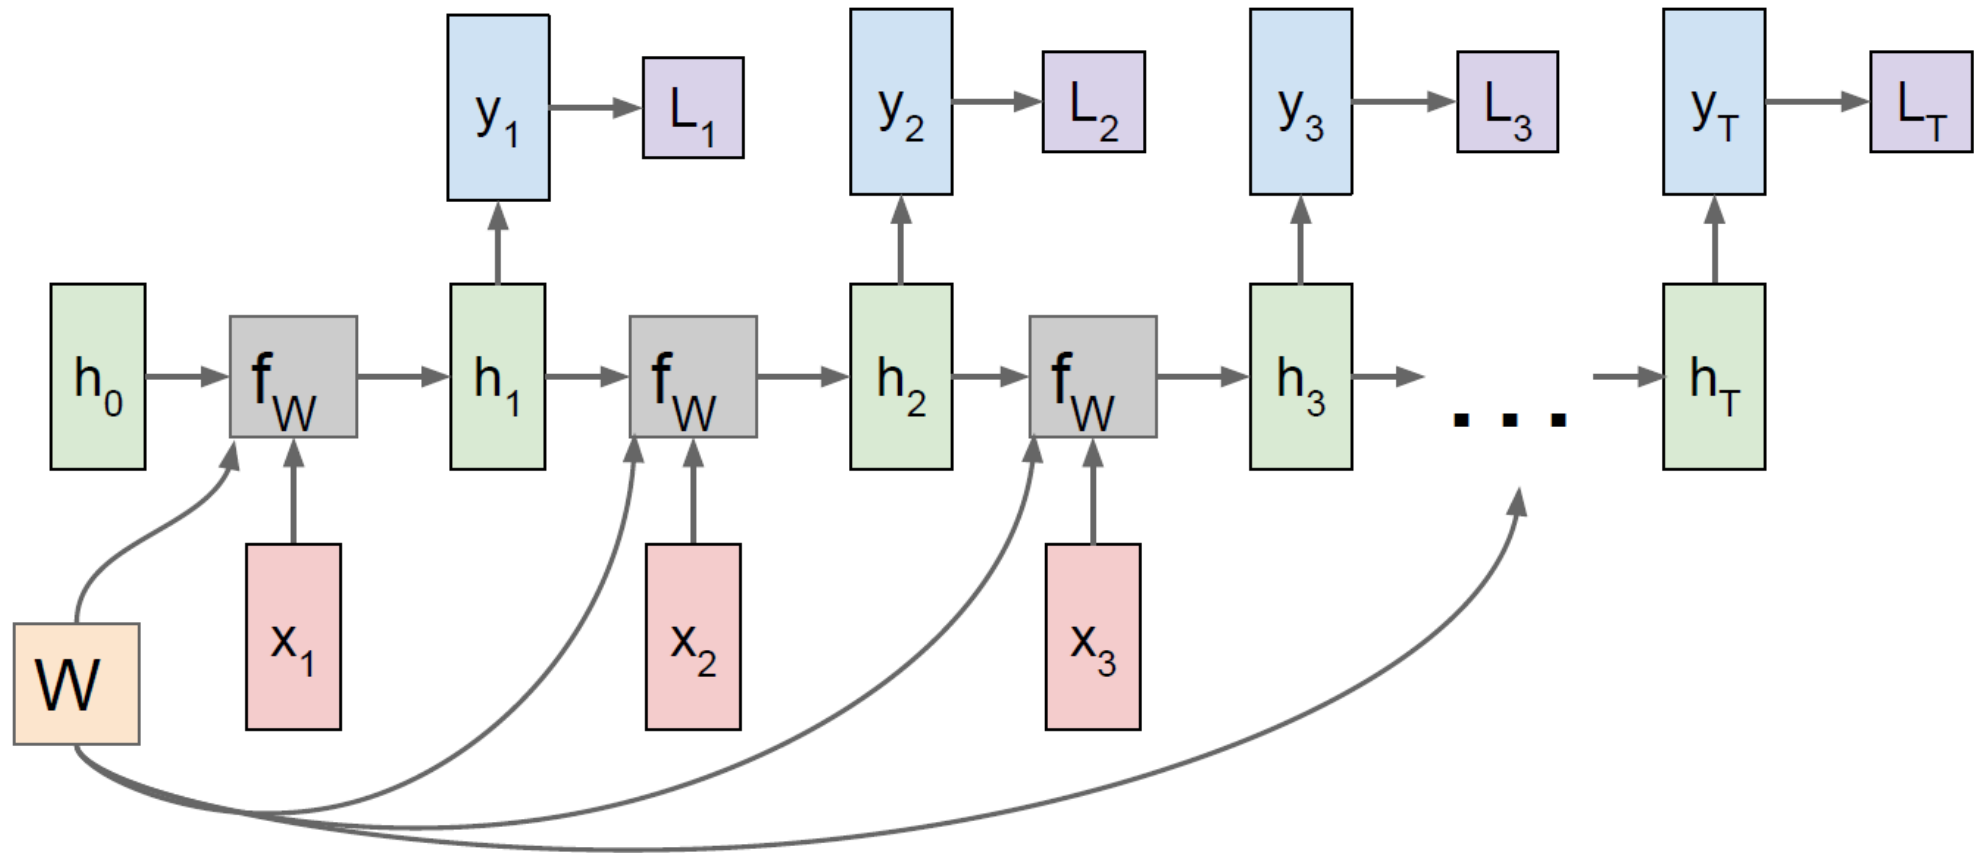
- Re-use the same weight matrix at every time-step



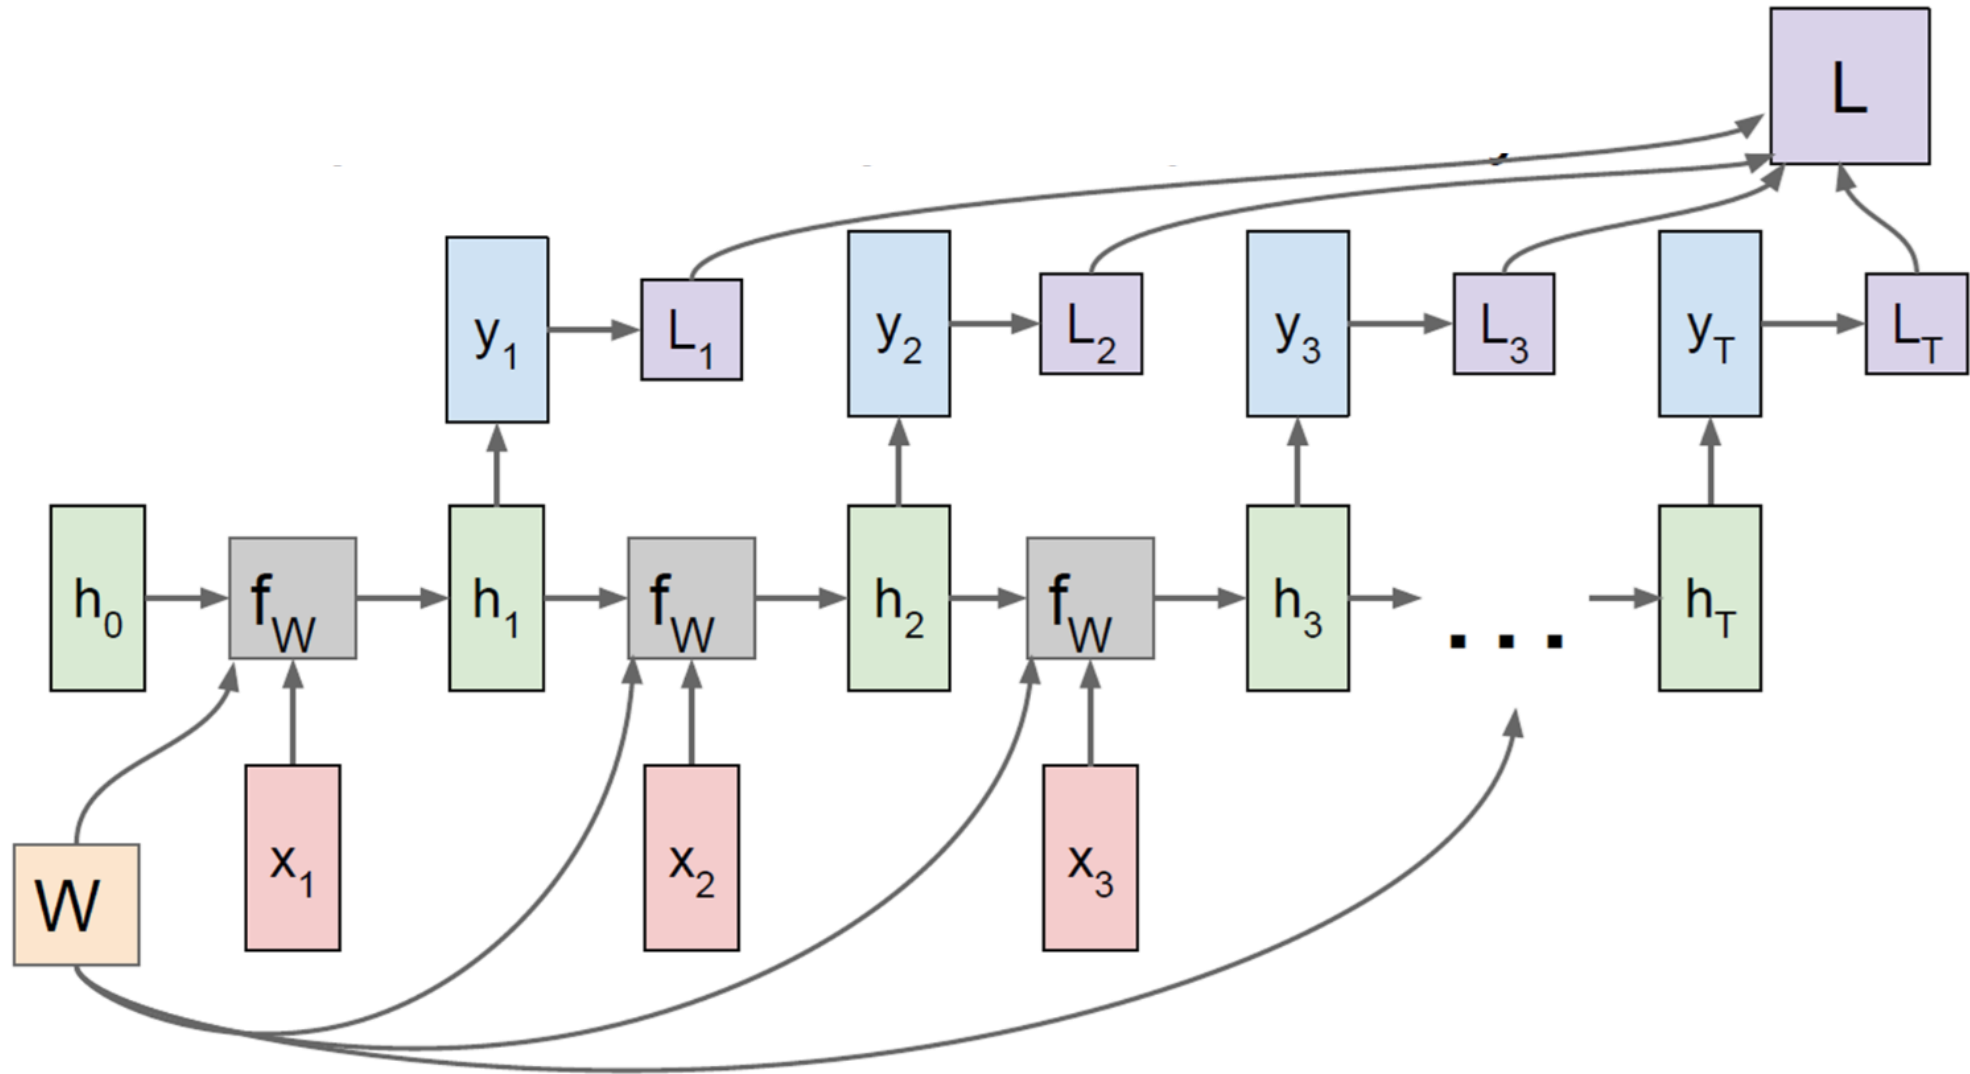
# Computational graph of RNN: many to many



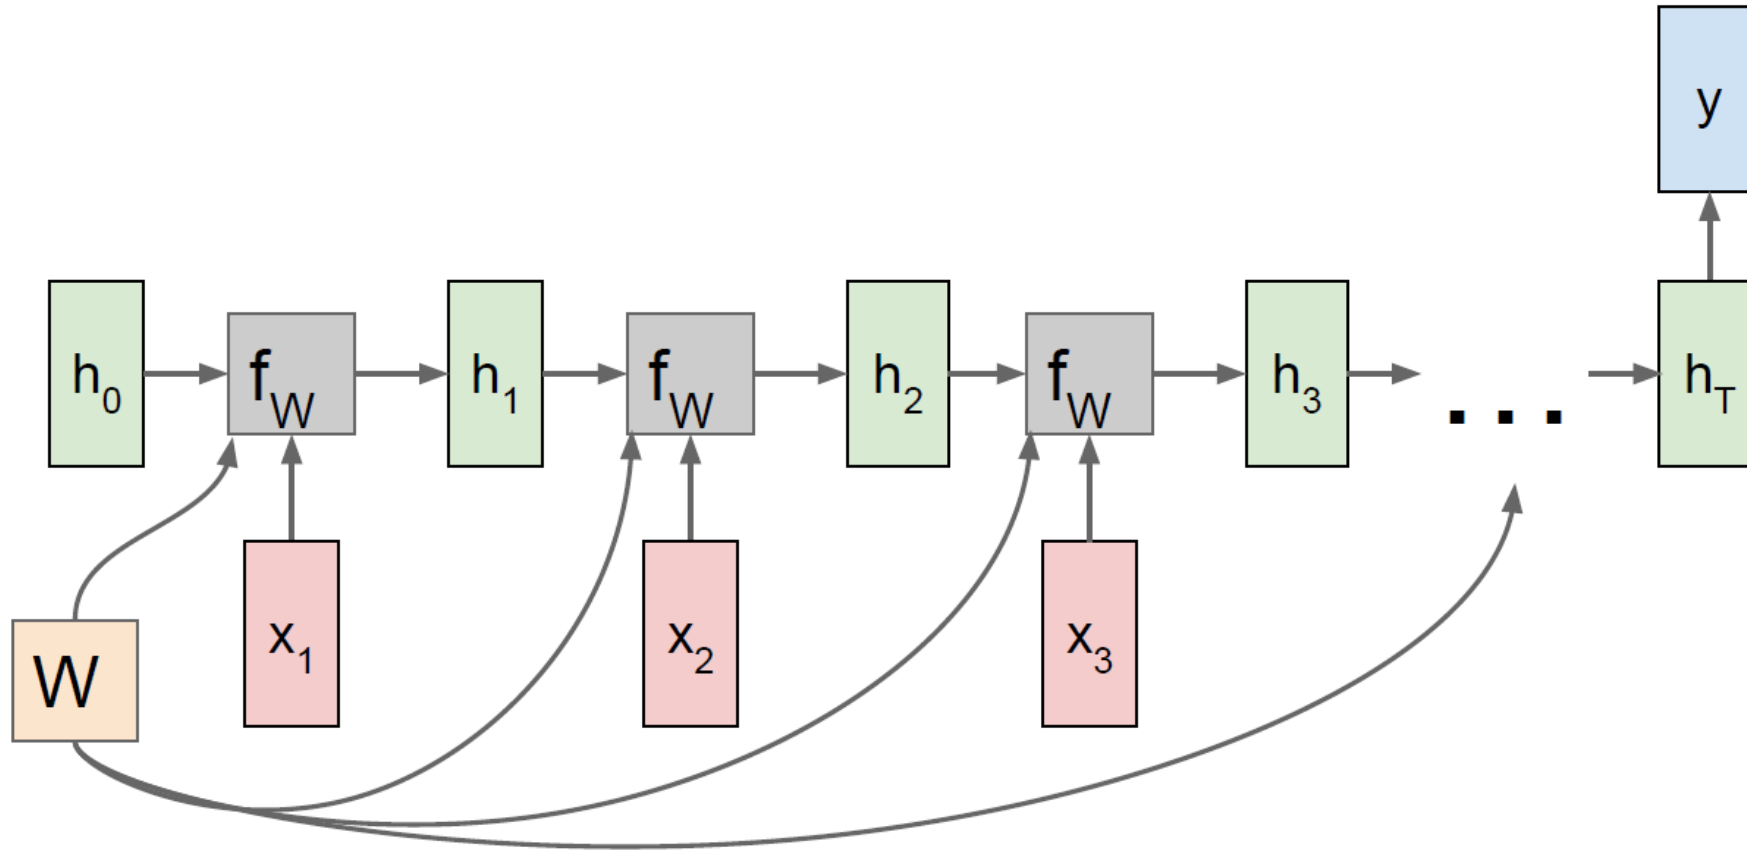
# Computational graph of RNN: many to many



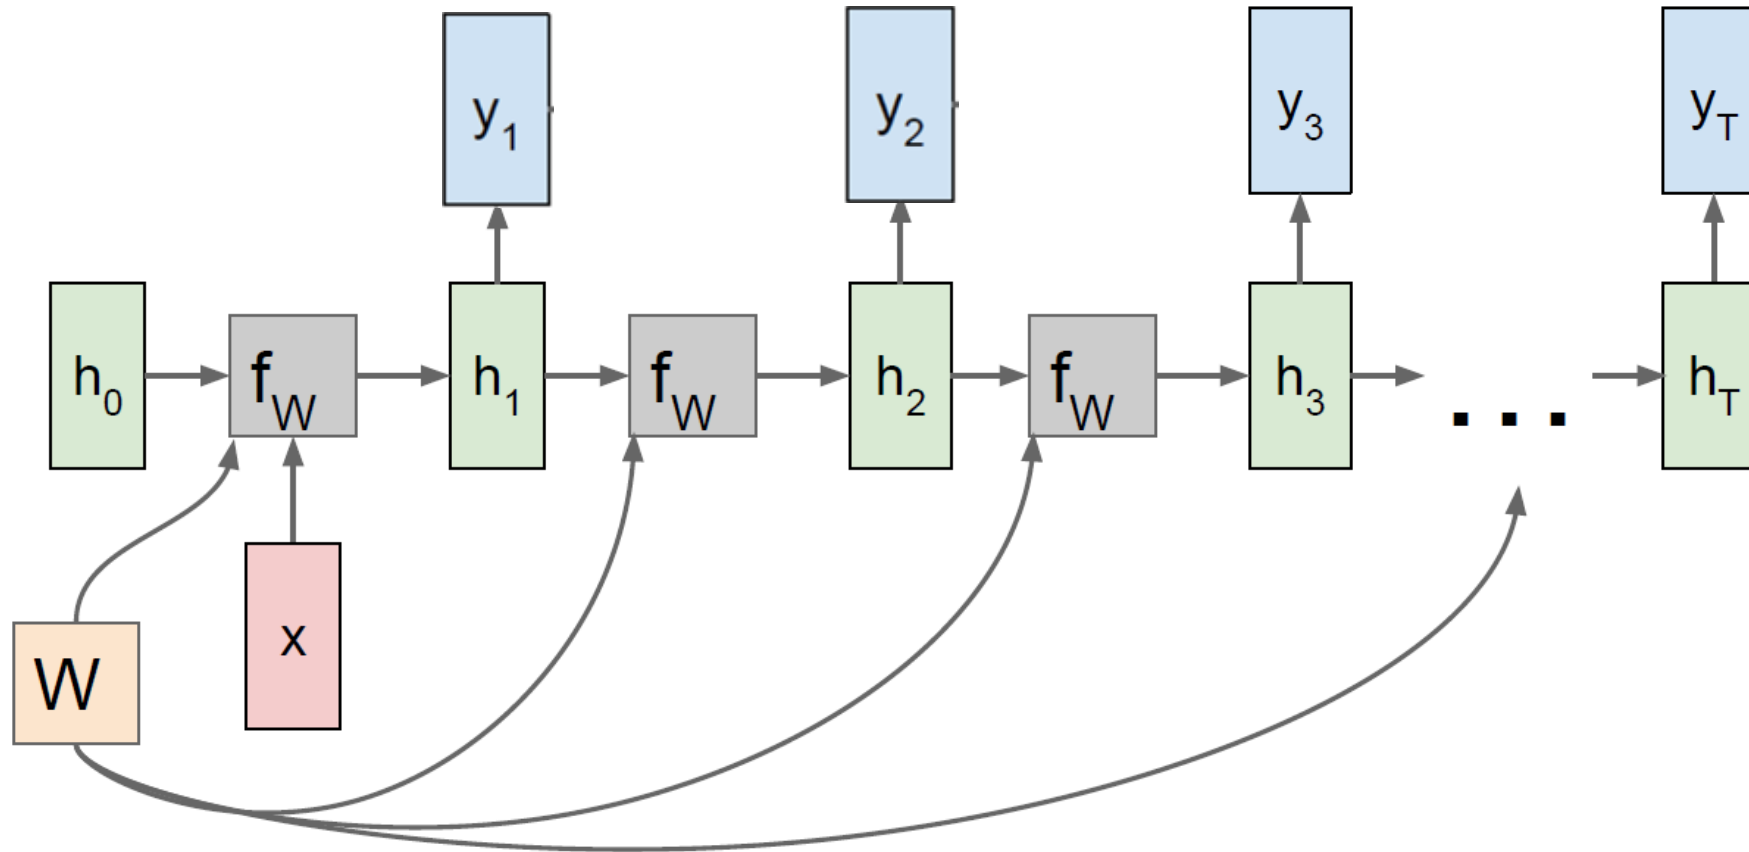
# Computational graph of RNN: many to many



# Computational graph of RNN: many to one



# Computational graph of RNN: one to many

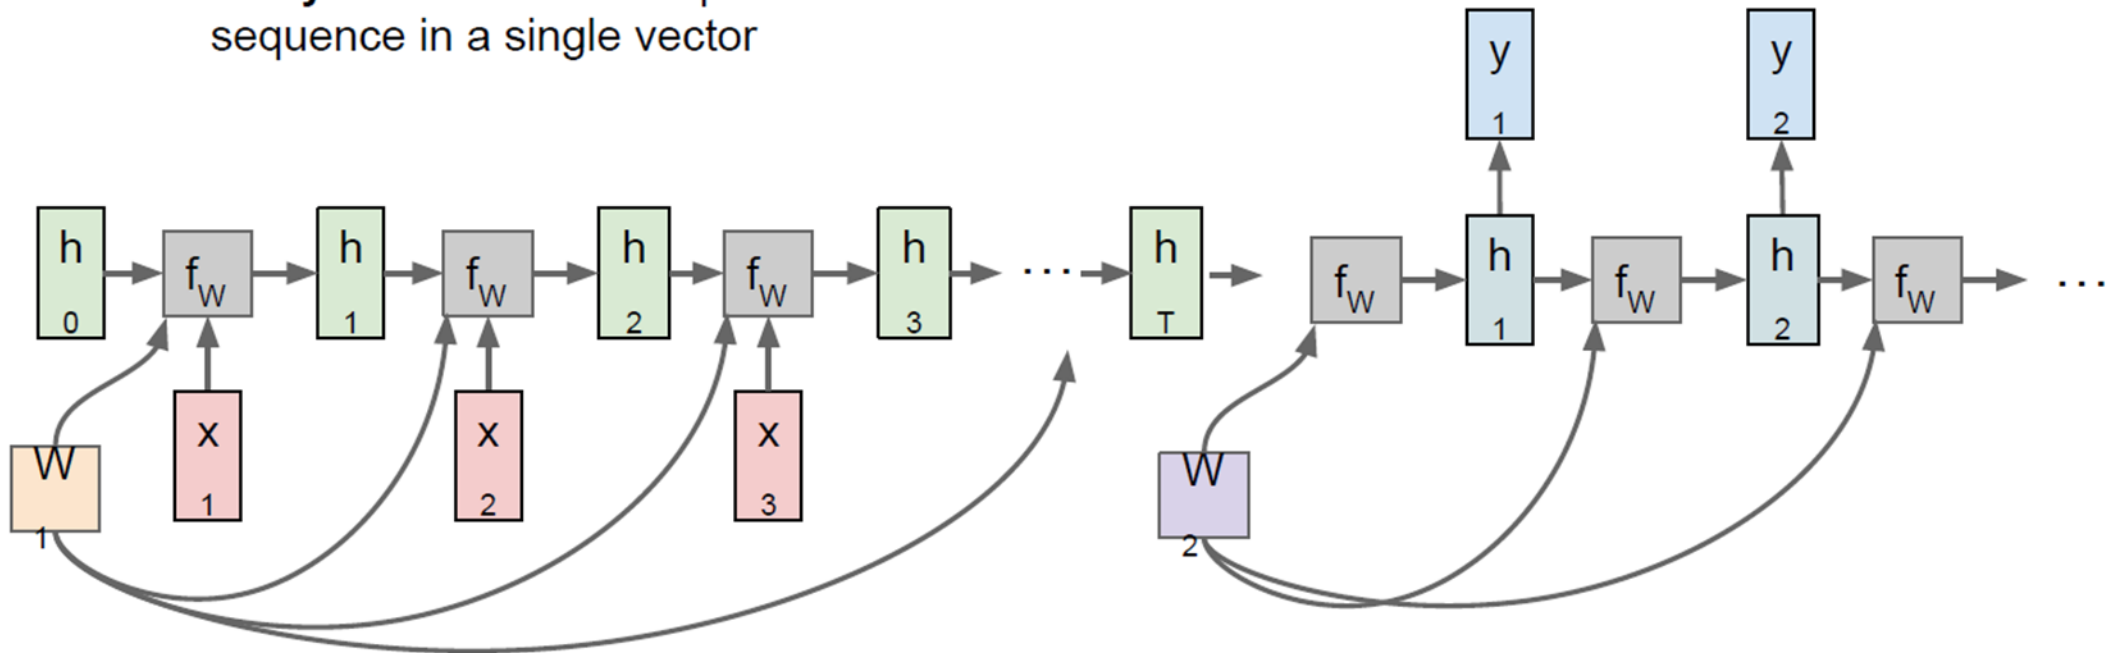


# Sequence to sequence

- Equals to man-to-one plus one-to-many

**Many to one:** Encode input sequence in a single vector


**One to many:** Produce output sequence from single input vector





# Different RNNs

one to one



Vanilla NN

one to many






Image captioning  
Text generation

many to one




Text classification  
Sentiment analysis

many to many



Machine translation  
Dialogue system


many to many



Stock price estimation  
Video frame classification

# Example: Character-level language model


- Given previous words/characters, predict the next
- Vocabulary: [h, e, l, o]
- Example of training sequence:  
“Hello”



# Example: Character-level language model

- Given previous words/characters, predict the next
- Vocabulary: [h, e, l, o]
- Example of training sequence:  
“Hello”

$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$




# Example: Character-level language model

- Given previous words/characters, predict the next
- Vocabulary: [h, e, l, o]
- Example of training sequence:  
“Hello”


$$h_t = \tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$

$$y_t = W_{hy}h_t$$



# Example: Character-level language model

- Given previous words/characters, predict the next
- Vocabulary: [h, e, l, o]
- Example of training sequence: "Hello"
- At test-time, sample characters one at a time, feed back to model



# Example of text generating


- Can we generate sonnets like a poet, using RNN?

## THE SONNETS

by William Shakespeare

From fairest creatures we desire increase,  
That thereby beauty's rose might never die,  
But as the ripper should by time decease,  
His tender heir might bear his memory:  
But thou, contracted to thine own bright eyes,  
Feed'st thy light's flame with self-substantial fuel,  
Making a famine where abundance lies,  
Thyself thy foe, to thy sweet self too cruel:  
Thou that art now the world's fresh ornament,  
And only herald to the gaudy spring,  
Within thine own bud buriest thy content,  
And tender churl mak'st waste in niggarding:  
Pity the world, or else this glutton be,  
To eat the world's due, by the grave and thee.

When forty winters shall besiege thy brow,  
And dig deep trenches in thy beauty's field,  
Thy youth's proud livery so gazed on now,  
Will be a tatter'd weed of small worth held:  
Then being asked, where all thy beauty lies,  
Where all the treasure of thy lusty days;  
To say, within thine own deep sunken eyes,  
Were an all-eating shame, and thriftless praise.  
How much more praise deserv'd thy beauty's use,  
If thou couldst answer 'This fair child of mine  
Shall sum my count, and make my old excuse,'  
Proving his beauty by succession thine!  
This were to be new made when thou art old,  
And see thy blood warm when thou feel'st it cold.



# Example result

at first:

tyntd-iafhatawiaoihrdemot lytdws e ,tfti, astai f ogoh eoase rrranbyne 'nhthnee e  
plia tklrqd t o idoe ns,smtt h ne etie h,hregtrs nigtike,aoaenns lng

↓ train more

"Tmont thithey" fomesscerliund  
Keushey. Thom here  
sheulke, anmerenith ol sivh I lalterthend Bleipile shuwy fil on aseterlome  
coaniogennc Phe lism thond hon at. MeiDimorotion in ther thize."

↓ train more

Aftair fall unsuch that the hall for Prince Velzonski's that me of  
her hearly, and behs to so arwage fiving were to it beloge, pavu say falling misfort  
how, and Gogition is so overelical and offer.

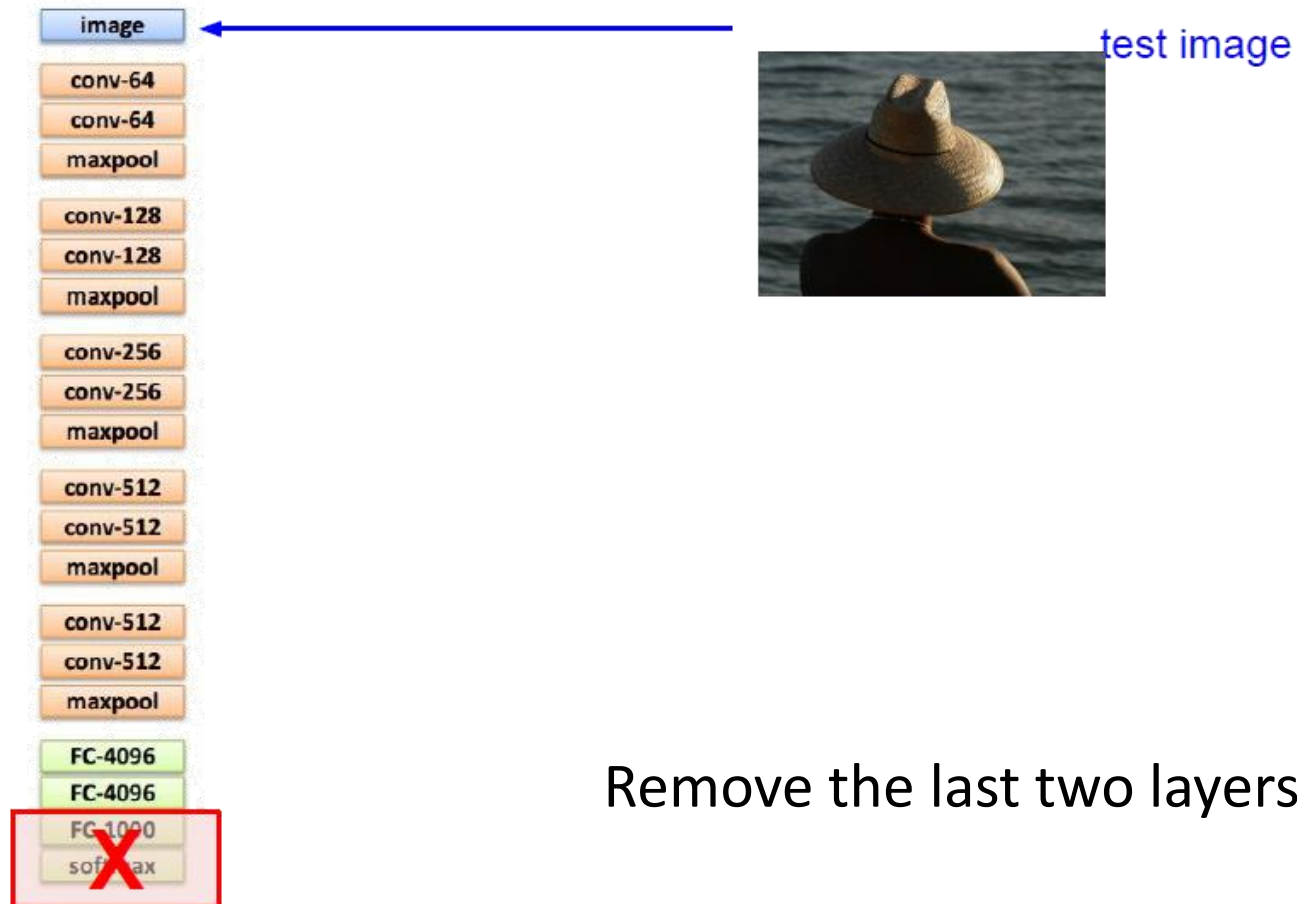
↓ train more

"Why do what that day," replied Natasha, and wishing to himself the fact the  
princess, Princess Mary was easier, fed in had oftended him.  
Pierre aking his soul came to the packs and drove up his father-in-law women.




# Image captioning

- Using RNN to explain the content in an Image



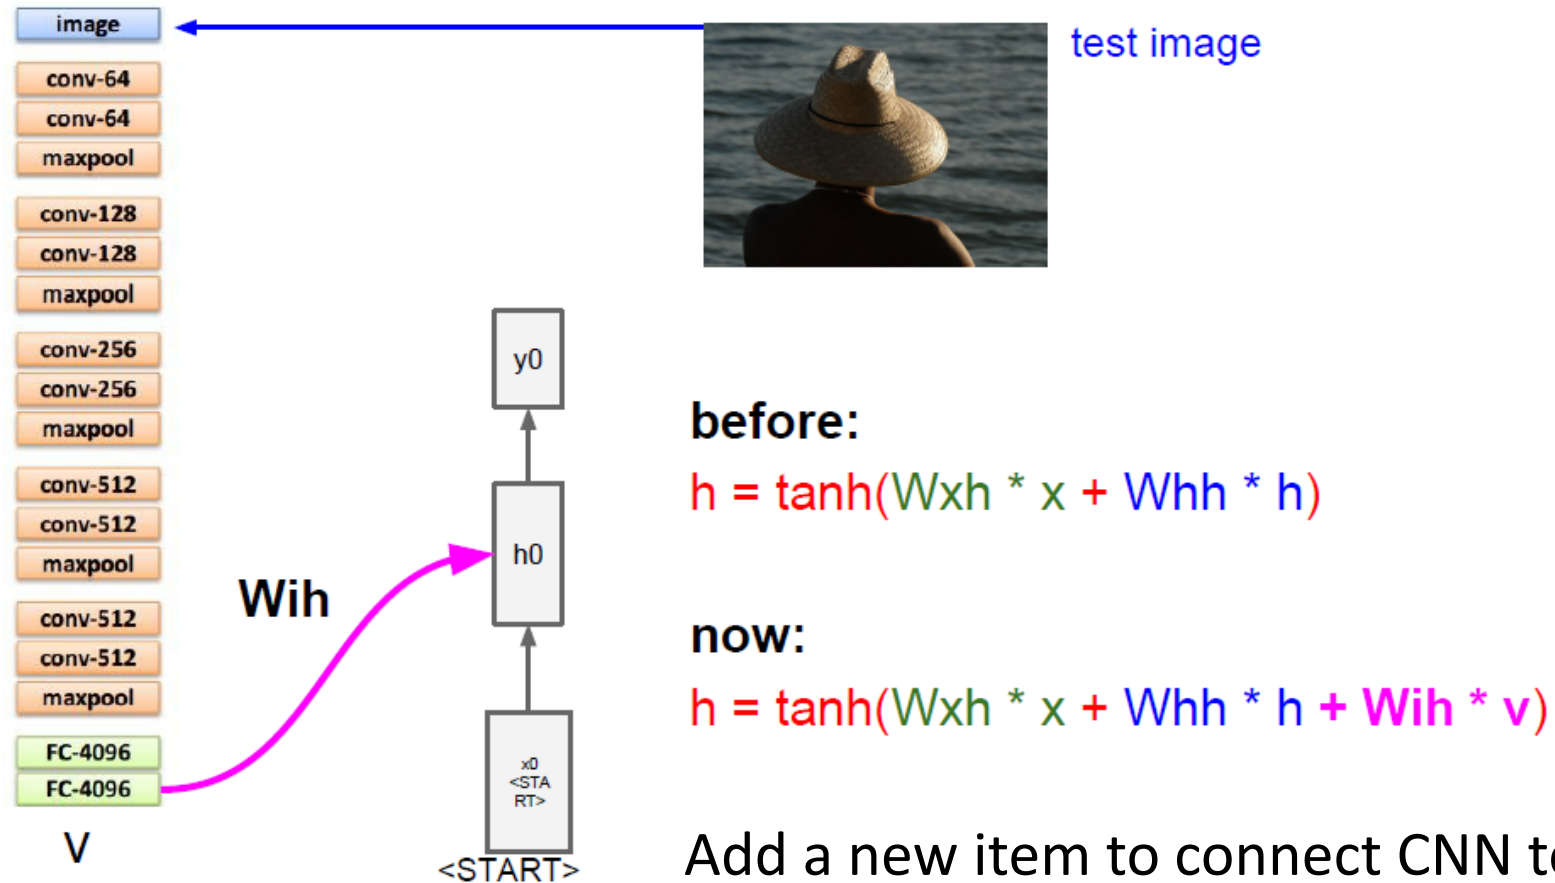


# Image captioning

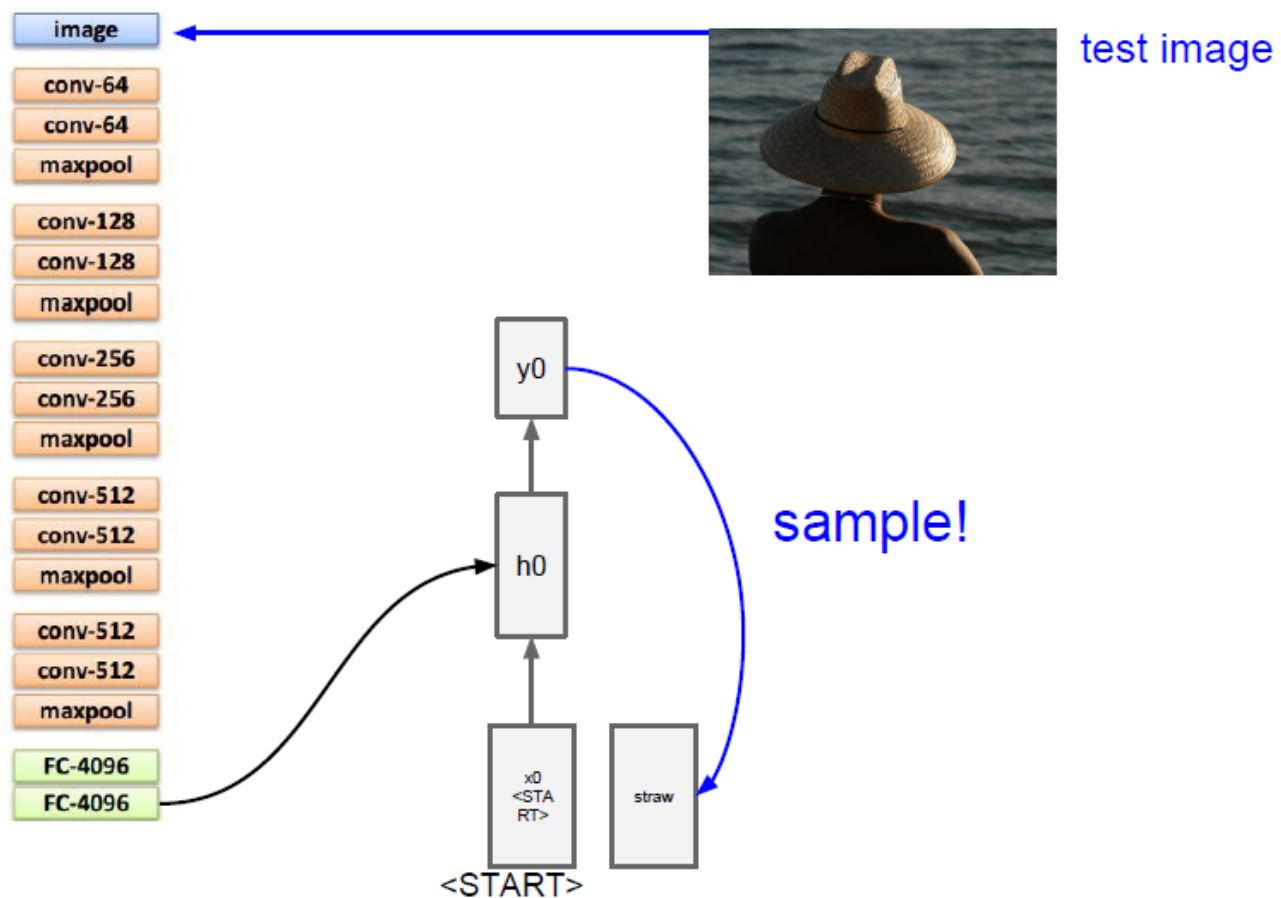


Preparing the initial input for RNN


# Image captioning



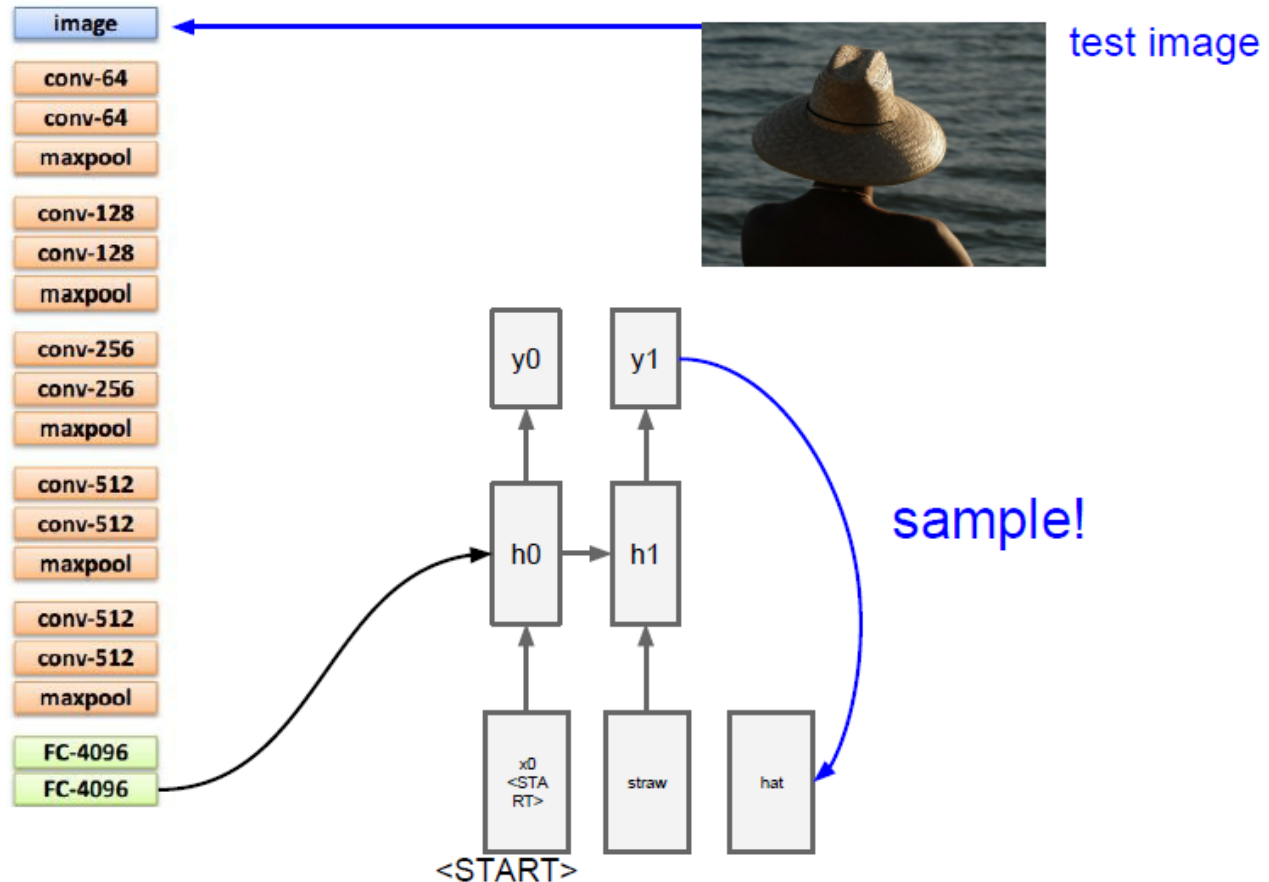
# Image captioning



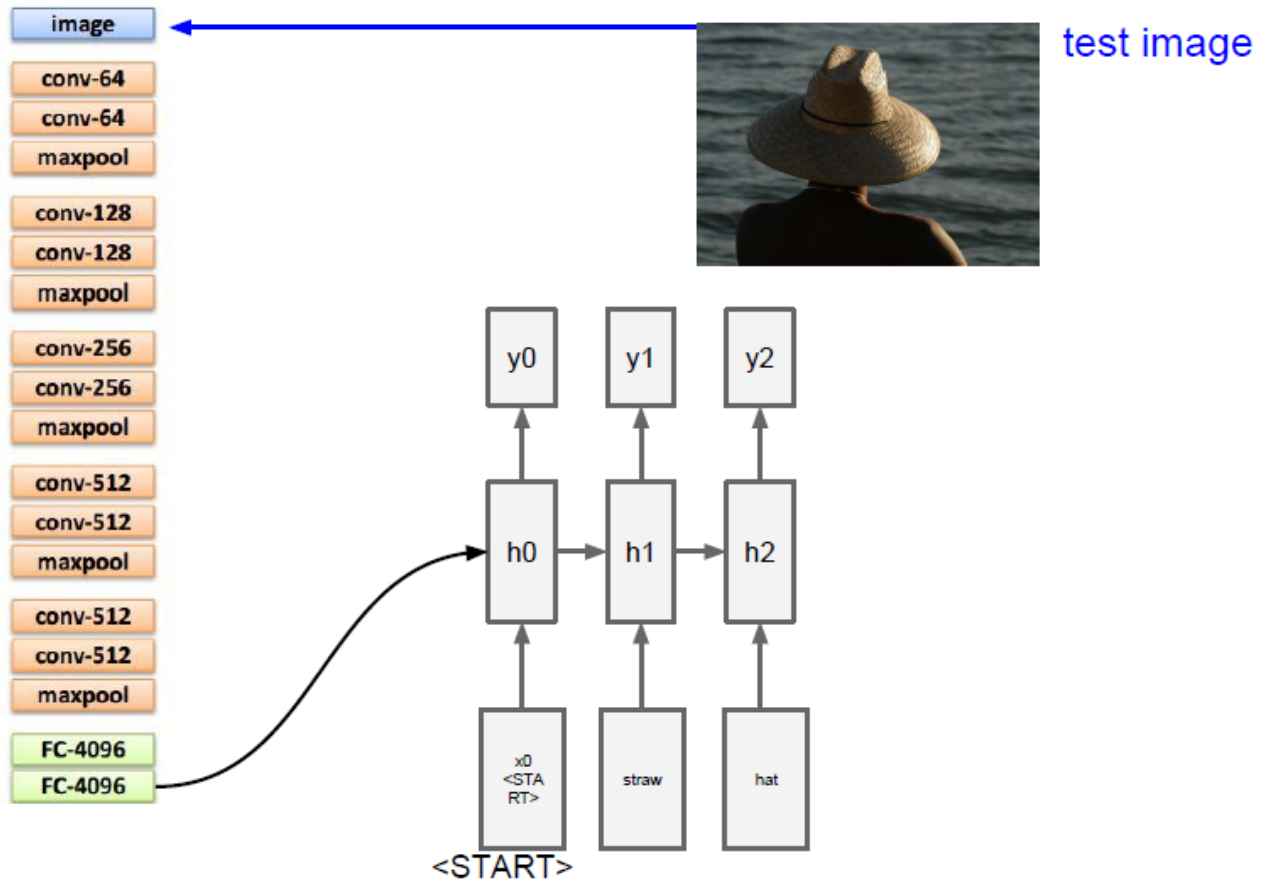
# Image captioning



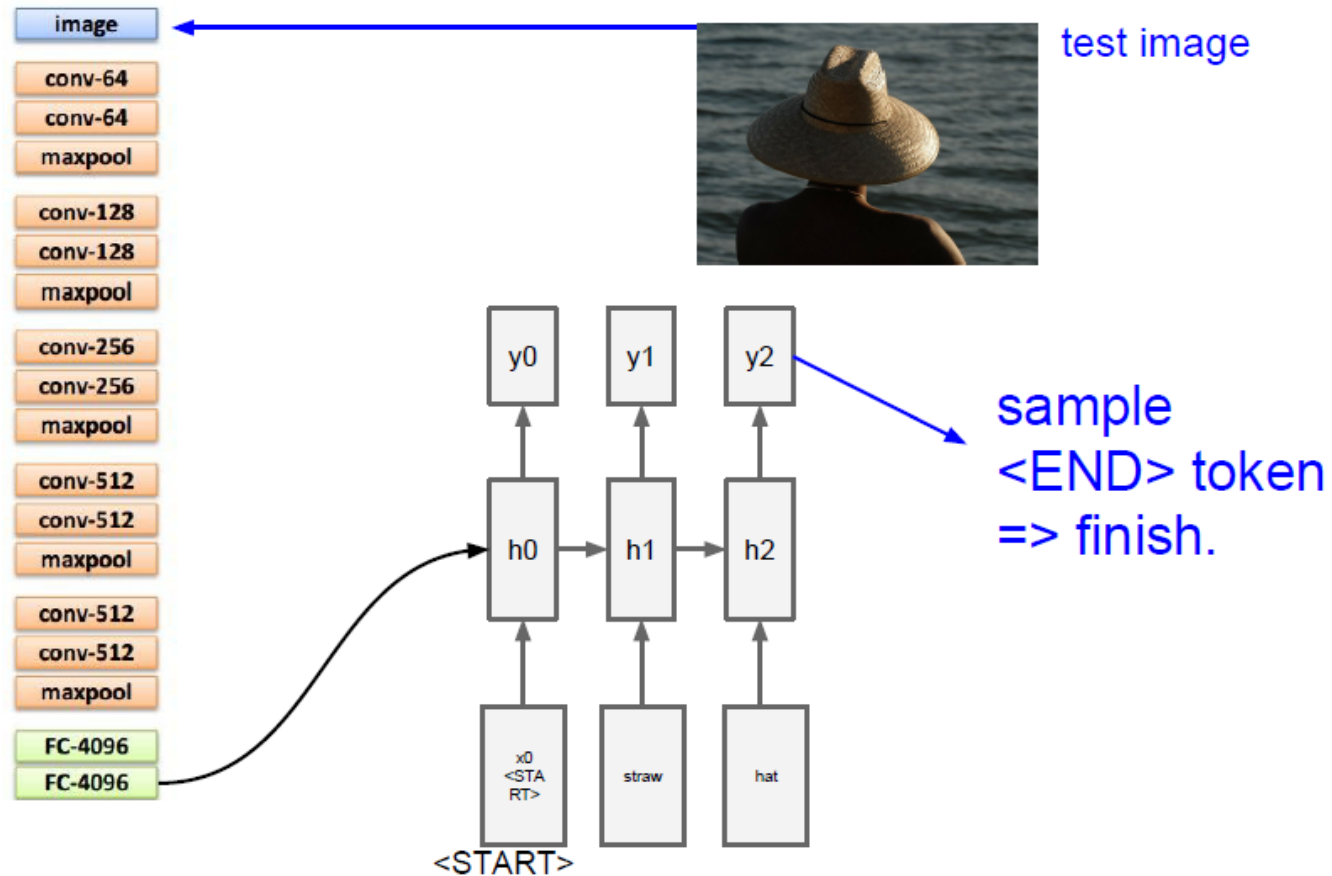
# Image captioning



# Image captioning



# Image captioning





# Image captioning: Example results



*A cat sitting on a suitcase on the floor*



*A cat is sitting on a tree branch*



*A dog is running in the grass with a frisbee*



*A white teddy bear sitting in the grass*



*Two people walking on the beach with surfboards*



*A tennis player in action on the court*



*Two giraffes standing in a grassy field*



*A man riding a dirt bike on a dirt track*




# Image captioning: Failure cases



*A woman is holding a cat in her hand*



*A person holding a computer mouse on a desk*



*A woman standing on a beach holding a surfboard*





*A bird is perched on a tree branch*



*A man in a baseball uniform throwing a ball*

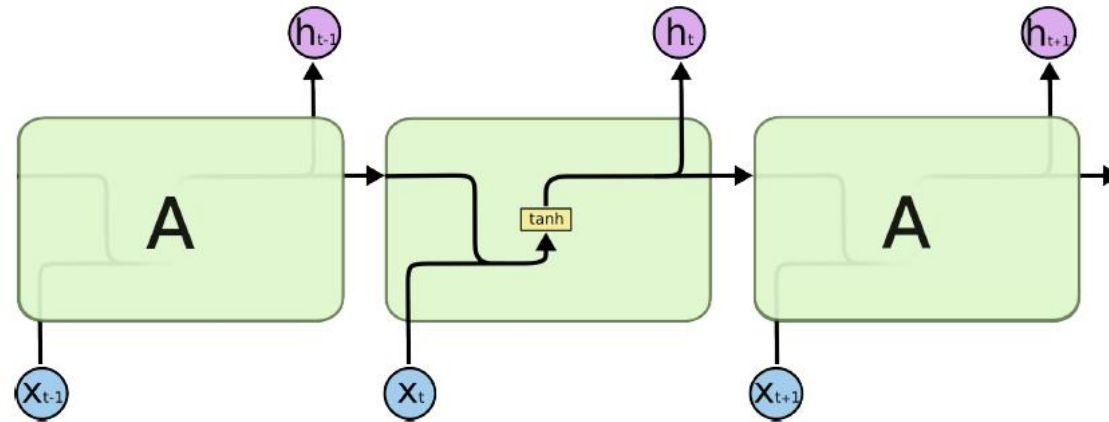
# The problem of long-term dependencies

- One of the appeals of RNNs is the idea that they might be able to connect previous information to the present task, such as using previous video frames might inform the understanding of the present frame
- RNN usually works well for recent information
  - E.g. The clouds are in the *sky*.
- But might work worse for further back context
  - E.g. I grew up in France... I speak fluent *French*.



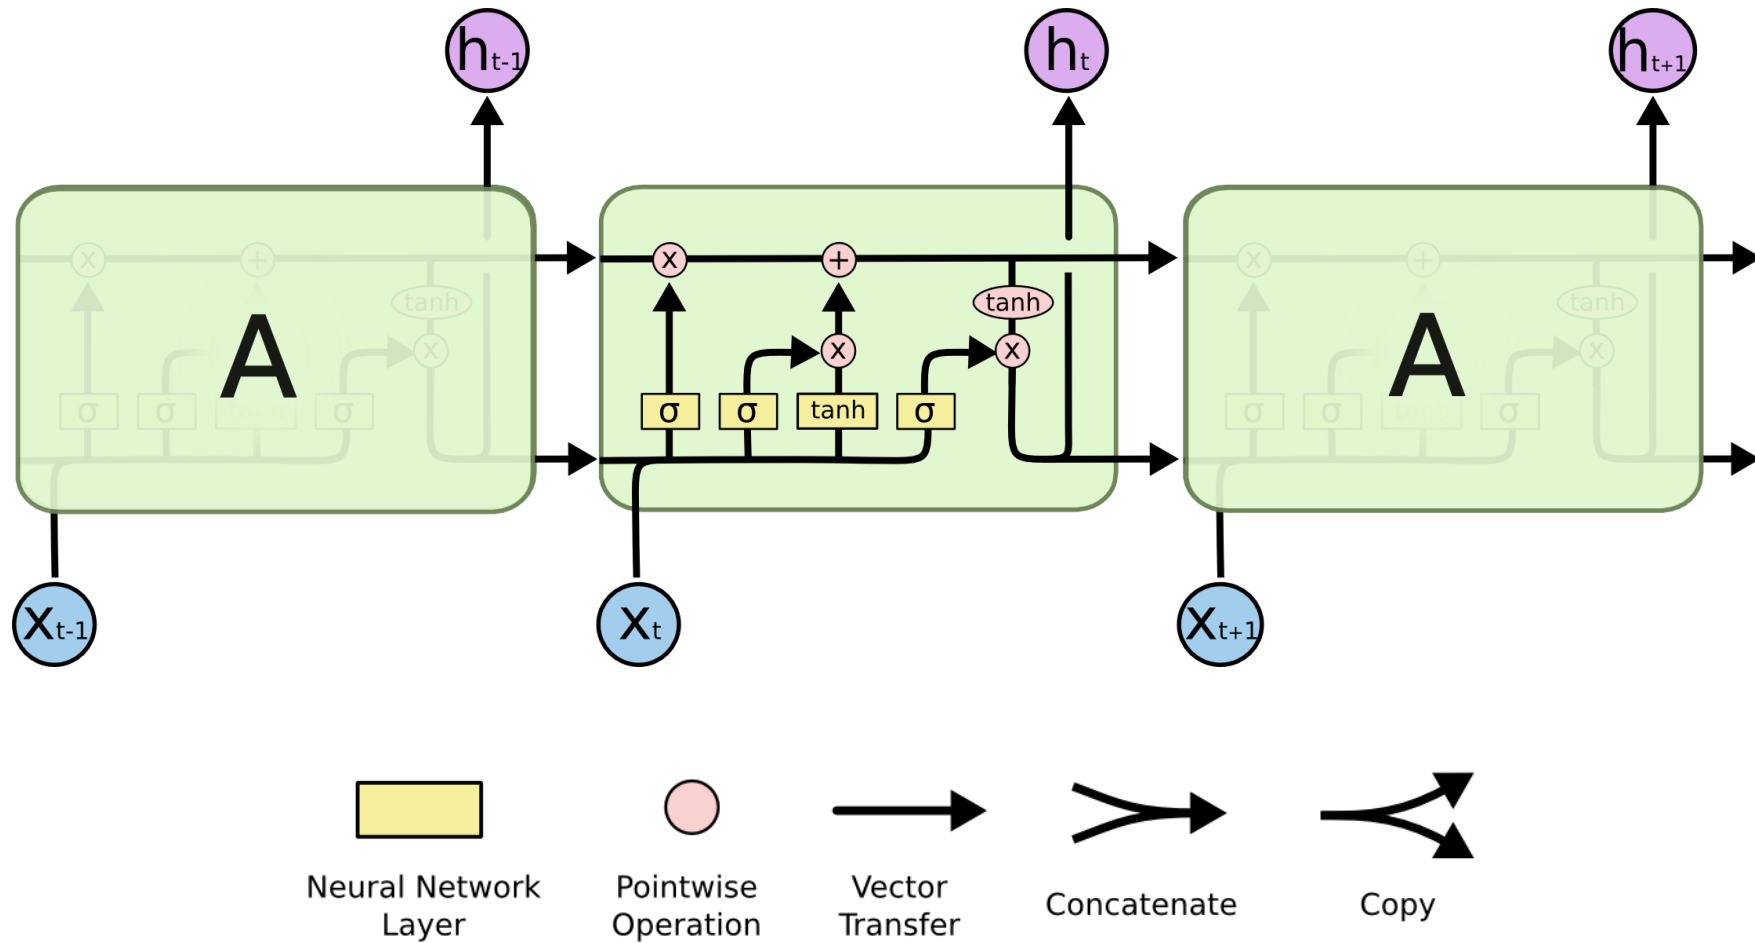
# Long short-term memory (LSTM)

- Hochreiter & Schmidhuber [1997]
- A variant of RNN, capable of learning long-term dependencies
- Can remember information for long periods of time



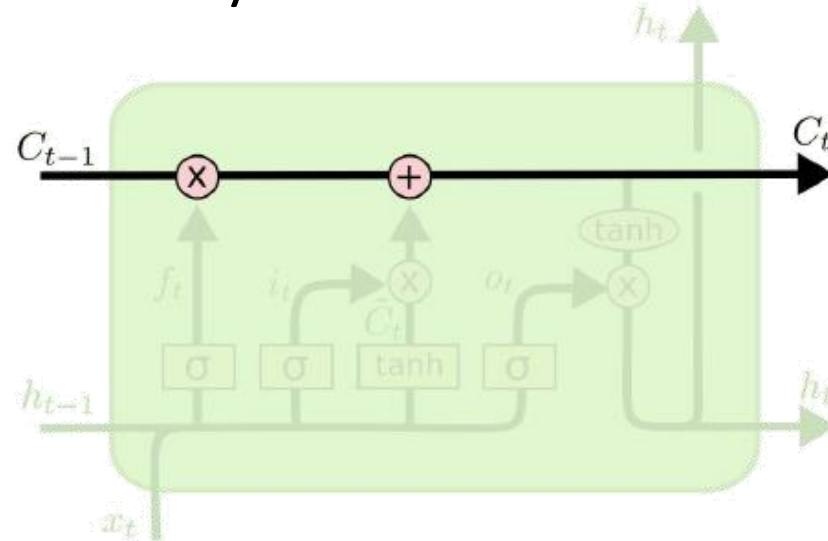
Remember the architecture of RNN cells, we will make many changes on it

# LSTM at a glance




# LSTM

- $C_t$  is the cell state
  - Horizontal line running through the top of the diagram
  - Like a conveyer
  - Run straight down the entire chain, with only some minor linear interactions
  - LSTM uses **gates** to carefully remove or add information to the cell state



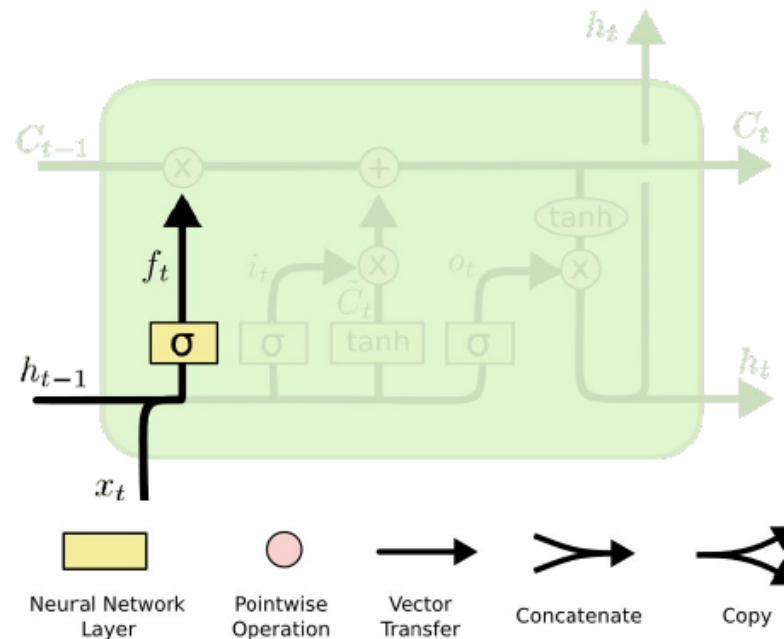
# LSTM (cont.)

- Three **gates** are governed by **sigmoid** units
- Describing how much of each component should be let through
- “let nothing through” (zero value)/“let everything through” (one value)



# LSTM (cont.)

- The first gate is **forget gate layer**
  - Decide what information we're going to throw away from the cell state
  - Output a number in (0,1) for each number in the cell state
  - 1/0: completely keep/remove this

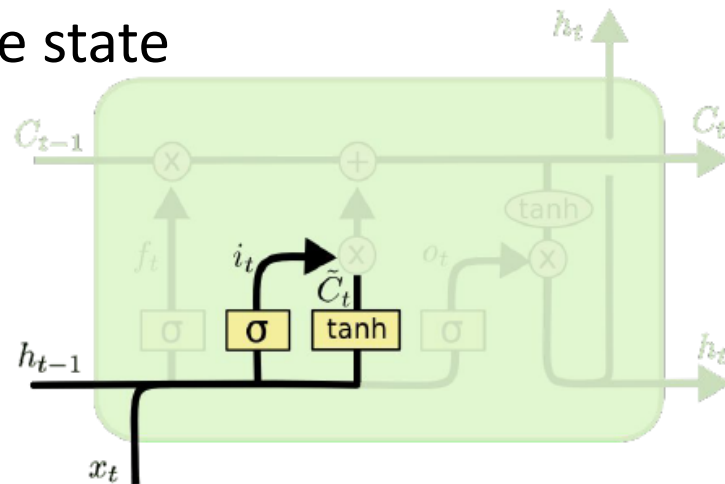


**Forget Gate:**

$$f_t = \sigma (W_f \cdot \underbrace{[h_{t-1}, x_t]}_{\text{Concatenate}} + b_f)$$

# LSTM (cont.)

- Next is to decide what new information we're going to store in the cell state
  - The sigmoid layer is the input gate layer, deciding which values we'll update
  - The tanh layer creates a vector of new candidate values  $\tilde{C}_t$  that should be added to the state



## Input Gate Layer

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

## New contribution to cell state


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

Classic neuron



# LSTM (cont.)

- Multiply the old state by  $f_t$ 
  - Forgetting the things from previous state
- Add new candidate values
- Update cell state

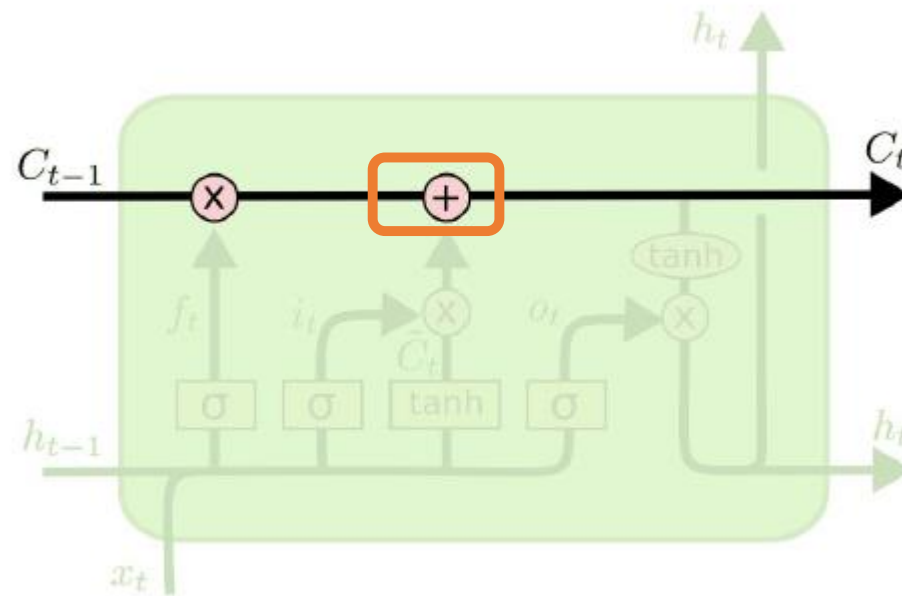


**Update Cell State (memory):**

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

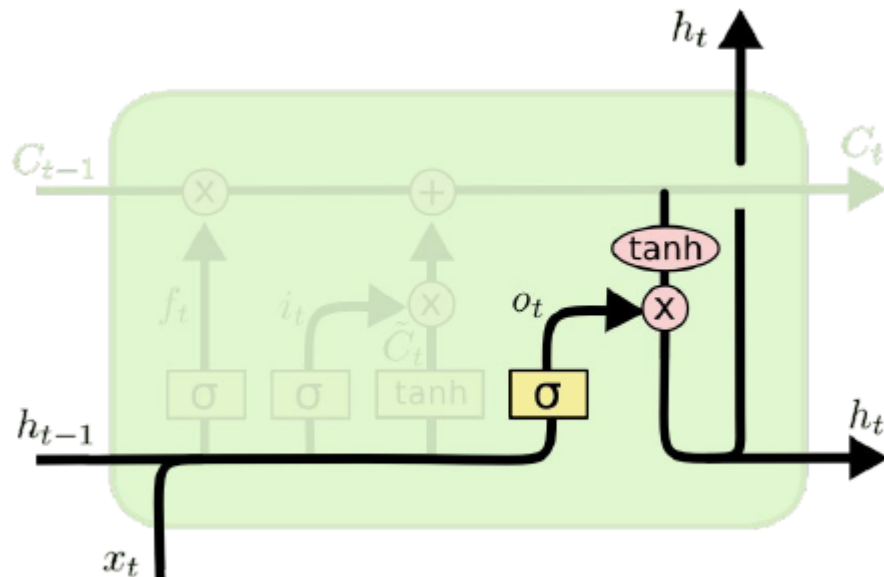
# LSTM (cont.)

- The **sum** operation avoids memory vanishing and exploding/vanishing backpropagation gradients, which are limitations of RNN



# LSTM (cont.)

- Last decide the output  $h_t$ 
  - Filtered (tanh) version of the cell state
  - The sigmoid layer decides what parts of the cell state we're going to output  
 $h_{t-1}, x_t$  are like *context*



## Output Gate Layer


$$o_t = \sigma(W_o [h_{t-1}, x_t] + b_o)$$

## Output to next layer

$$h_t = o_t * \tanh(C_t)$$

# Gated recurrent unit (GRU)

- GRU is a variants of LSTM
  - Combines the forget and input gates into a single “update gate”
  - Also merges the cell state and hidden state



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# LSTM hyperparameter tuning

- Watch out for *overfitting*
- Regularization helps: regularization methods include L1, L2, and dropout among others.
- The larger the network, the more powerful, but also easier to overfit. Don't want to try to learn a million parameters from 10,000 examples
  - parameters > examples = trouble
- More data is usually better
- Train over multiple epochs (complete passes through the dataset)
- The learning rate is the single most important hyper parameter